

# Address Translation

## Introduction

Hammurabi Mendes  
Fall 18

# Paging

# Paging

- Both *virtual* and *physical* address spaces are divided in “buckets”

# Paging

- Both *virtual* and *physical* address spaces are divided in “buckets”
  - In the virtual space, they’re called *pages*

# Paging

- Both *virtual* and *physical* address spaces are divided in “buckets”
  - In the virtual space, they’re called *pages*
  - In the physical space, they’re called *frames*

# Paging

- Both *virtual* and *physical* address spaces are divided in “buckets”
  - In the virtual space, they’re called *pages*
  - In the physical space, they’re called *frames*
- Pages and frames have the *same, fixed* size

# Paging

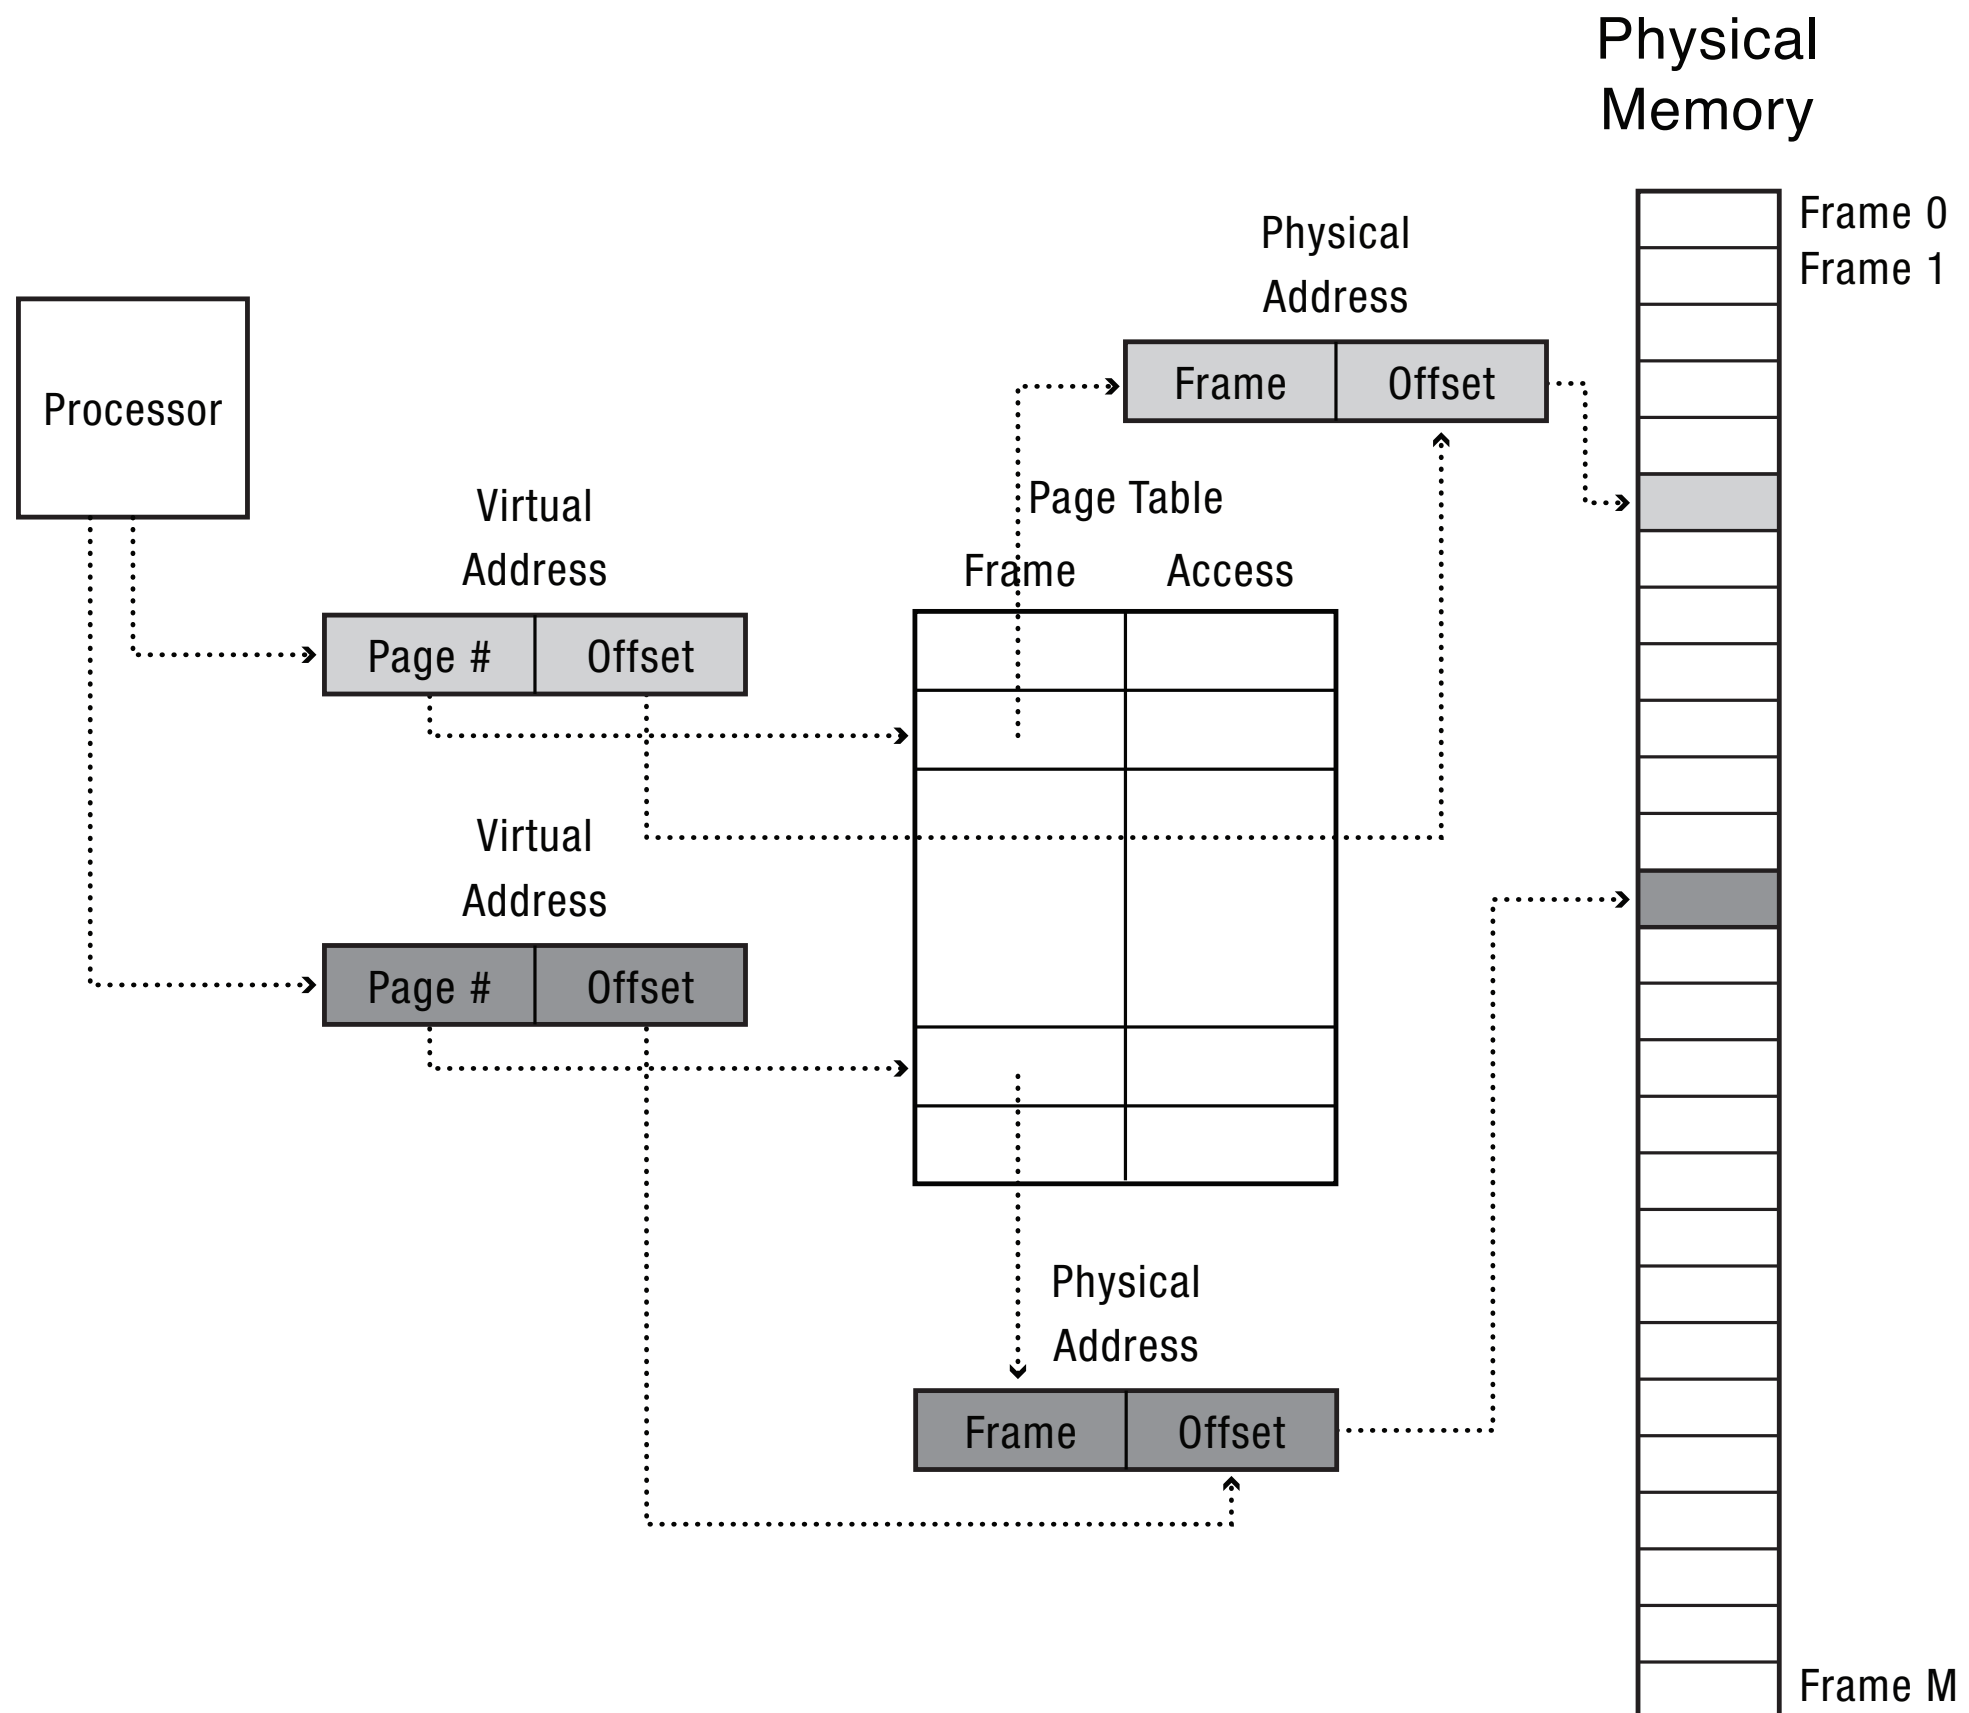
- Both *virtual* and *physical* address spaces are divided in “buckets”
  - In the virtual space, they’re called *pages*
  - In the physical space, they’re called *frames*
- Pages and frames have the *same, fixed* size
  - A power of 2: say  $2^{12} = 4096$  B

# Paging

- Both *virtual* and *physical* address spaces are divided in “buckets”
  - In the virtual space, they’re called *pages*
  - In the physical space, they’re called *frames*
- Pages and frames have the *same, fixed* size
  - A power of 2: say  $2^{12} = 4096$  B
  - This prevents external fragmentation!



# Paging



# Paging

x86-64 gives you a 48-bit virtual address space:

Virtual  
Address

Page #	Offset
--------	--------

# Paging

x86-64 gives you a 48-bit virtual address space:

Virtual  
Address



12 bits

# Paging

x86-64 gives you a 48-bit virtual address space:

Virtual  
Address

Page #	Offset
--------	--------

12 bits

Page size:  
 $2^{12} = 4096$  B

# Paging

x86-64 gives you a 48-bit virtual address space:

Virtual  
Address

Page #	Offset
--------	--------

36 bits    12 bits

Page size:  
 $2^{12} = 4096$  B

# Paging

x86-64 gives you a 48-bit virtual address space:

Virtual  
Address

Page #	Offset
--------	--------

36 bits    12 bits

Page size:  
 $2^{12} = 4096 \text{ B}$

Maximum process image:  
 $2^{48} = 281474976710655 \text{ B} = 256 \text{ TB}$

# Paging

# Paging

- Advantages over Segmentation



# Paging

- Advantages over Segmentation
  - No external fragmentation
    - Only *internal fragmentation*: last page is not fully used

# Paging

- Advantages over Segmentation
  - No external fragmentation
    - Only *internal fragmentation*: last page is not fully used
  - Can load code/data on demand as the process executes

# Paging

- Advantages over Segmentation
  - No external fragmentation
    - Only *internal fragmentation*: last page is not fully used
  - Can load code/data on demand as the process executes
    - When the program references a non-mapped virtual address, the hardware generates a *page fault (interrupt)*

# Paging

- Advantages over Segmentation
  - No external fragmentation
    - Only *internal fragmentation*: last page is not fully used
  - Can load code/data on demand as the process executes
    - When the program references a non-mapped virtual address, the hardware generates a *page fault (interrupt)*
    - The OS loads the page and resumes the program

# Paging

- Advantages over Segmentation
  - No external fragmentation
    - Only *internal fragmentation*: last page is not fully used
  - Can load code/data on demand as the process executes
    - When the program references a non-mapped virtual address, the hardware generates a *page fault (interrupt)*
    - The OS loads the page and resumes the program
  - Out of memory? Store unused pages in the disk!

# Paging

# Paging

- Drawbacks:

# Paging

- Drawbacks:
  - Requires more space to translate address



# Paging

- Drawbacks:
  - Requires more space to translate address
    - *Naively*, we need an entry for every page!

# Paging


- Drawbacks:
  - Requires more space to translate address
    - *Naively*, we need an entry for every page!
    - That's 262144 entries for a 1GB space

# Paging

- Drawbacks:
  - Requires more space to translate address
    - *Naively*, we need an entry for every page!
    - That's 262144 entries for a 1GB space
  - Keeping track of free/used physical frames


# Paging

- Drawbacks:
  - Requires more space to translate address
    - *Naively*, we need an entry for every page!
    - That's 262144 entries for a 1GB space
  - Keeping track of free/used physical frames
    - *Simplistically*, a bitmap solves the problem



# Paging

- Drawbacks:
  - Requires more space to translate address
    - *Naively*, we need an entry for every page!
    - That's 262144 entries for a 1GB space
  - Keeping track of free/used physical frames
    - *Simplistically*, a bitmap solves the problem




1)

Thanks Wikipedia user "Didia"

# Paging

- Drawbacks:
  - Requires more space to translate address
    - *Naively*, we need an entry for every page!
    - That's 262144 entries for a 1GB space
  - Keeping track of free/used physical frames
    - *Simplistically*, a bitmap solves the problem




1)  
2)

Thanks Wikipedia user "Didia"  
This means "allocation units"

# Paging

- Drawbacks:
  - Requires more space to translate address
    - *Naively*, we need an entry for every page!
    - That's 262144 entries for a 1GB space
  - Keeping track of free/used physical frames
    - *Simplistically*, a bitmap solves the problem




1)  
2)  
3)

Thanks Wikipedia user "Didia"  
This means "allocation units"  
Thanks Google Translate

# Paging

- Drawbacks:
  - Requires more space to translate address
    - *Naively*, we need an entry for every page!
    - That's 262144 entries for a 1GB space
  - Keeping track of free/used physical frames
    - *Simplistically*, a bitmap solves the problem





*Of course*

we will use better data structures  
to tackle both problems

# Multilevel Paging

# Multilevel Paging

- Never heard of a program using 256TB of memory

# Multilevel Paging

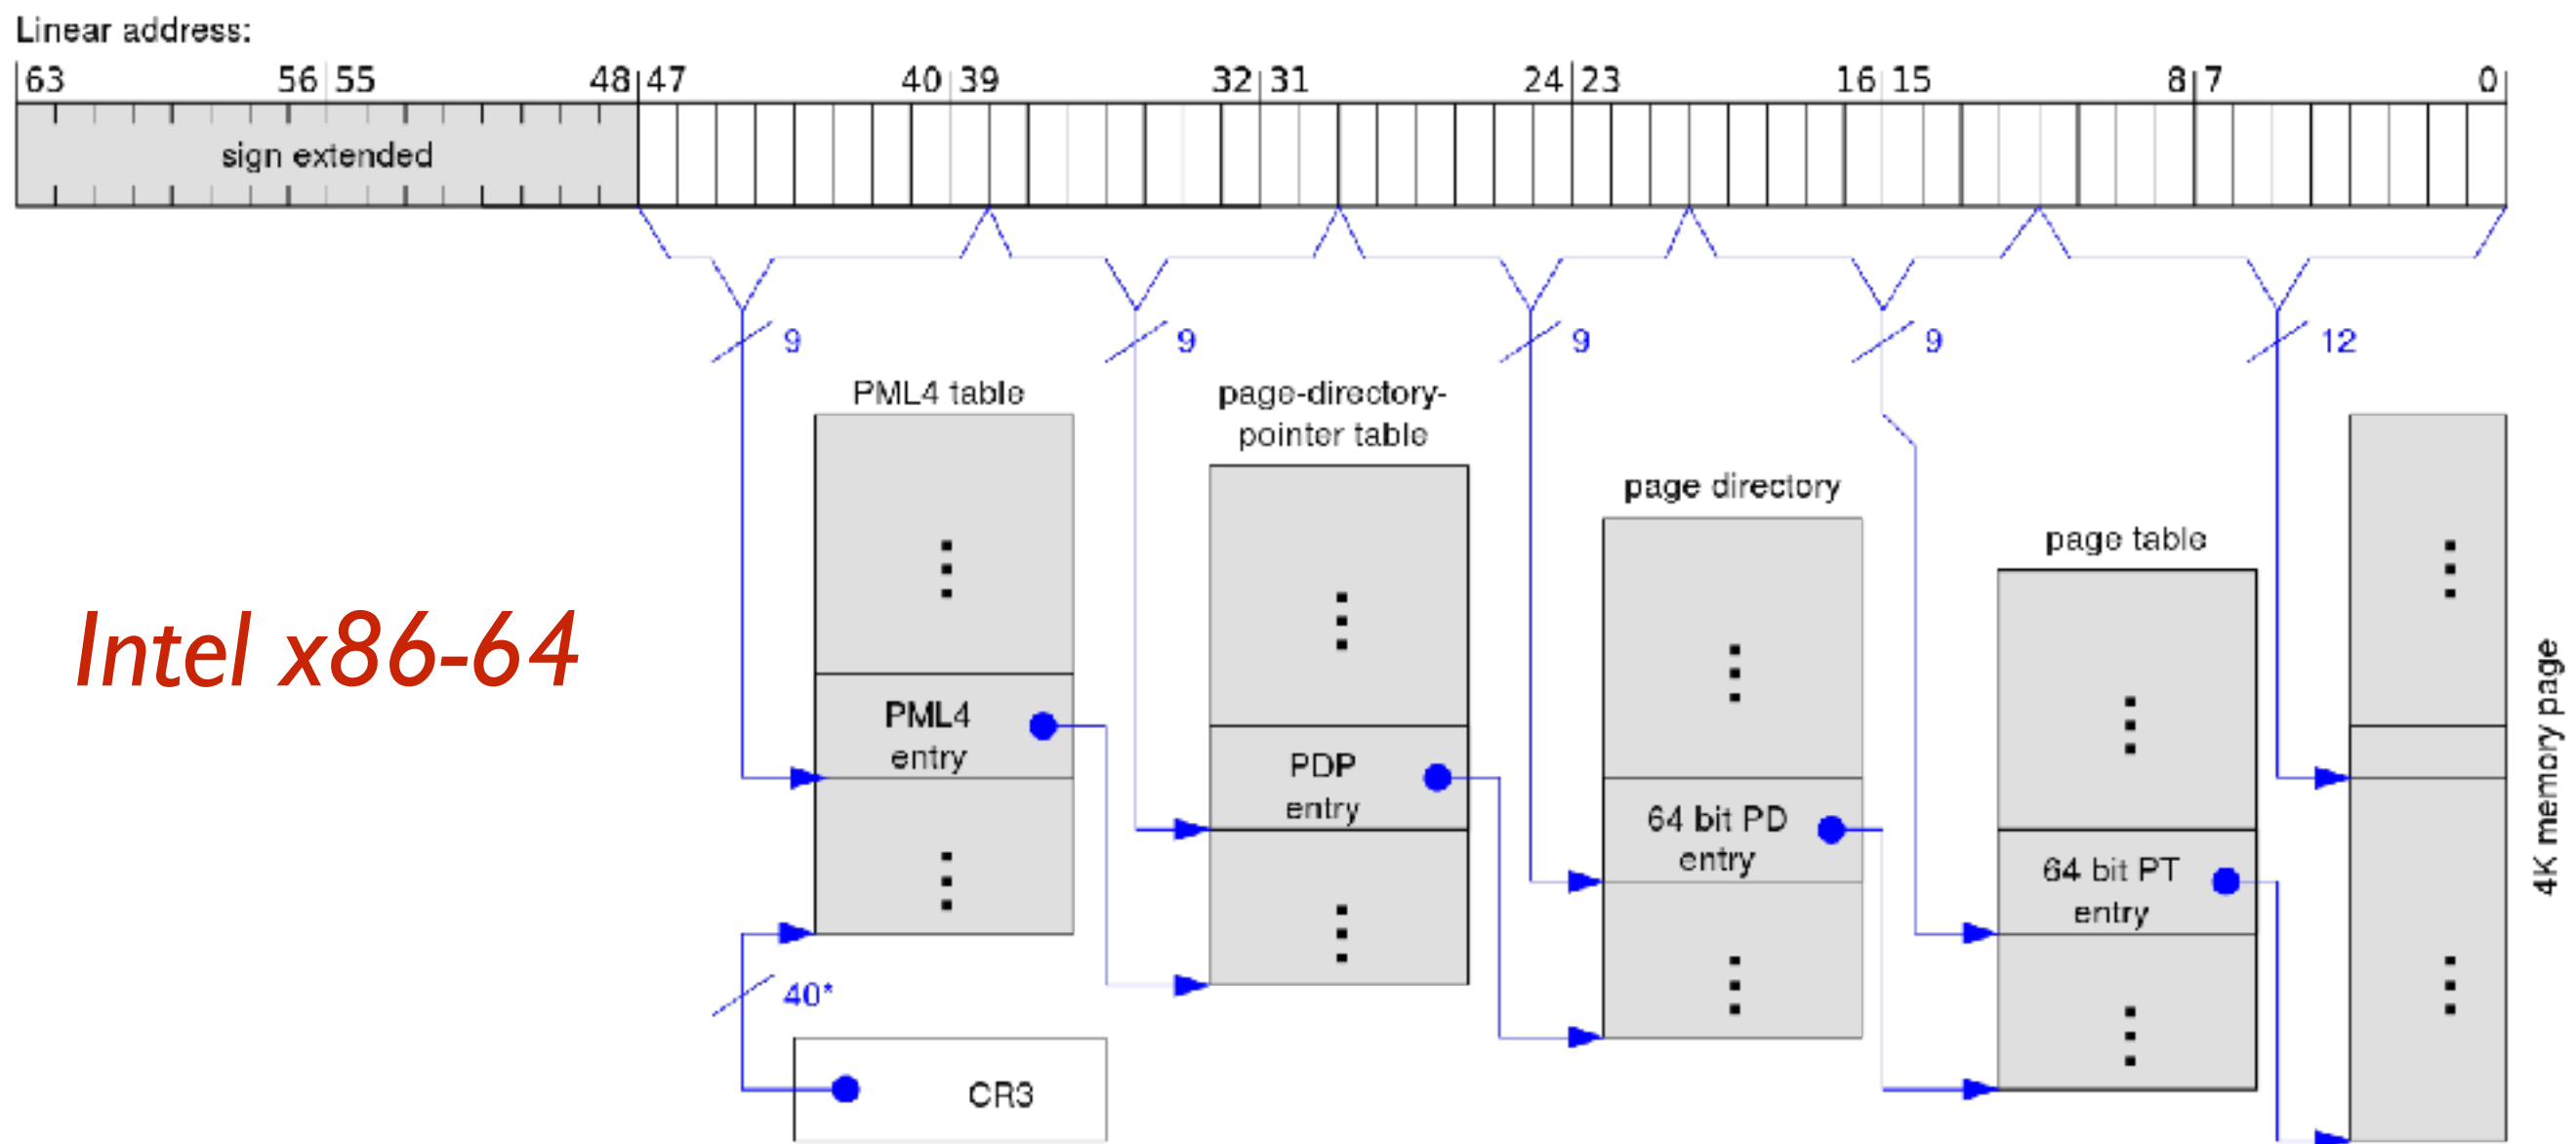
- Never heard of a program using 256TB of memory
  - Page table is a sparse data structure

# Multilevel Paging

- Never heard of a program using 256TB of memory
  - Page table is a sparse data structure
  - Should be divided, with pieces allocated on demand!

# Multilevel Paging

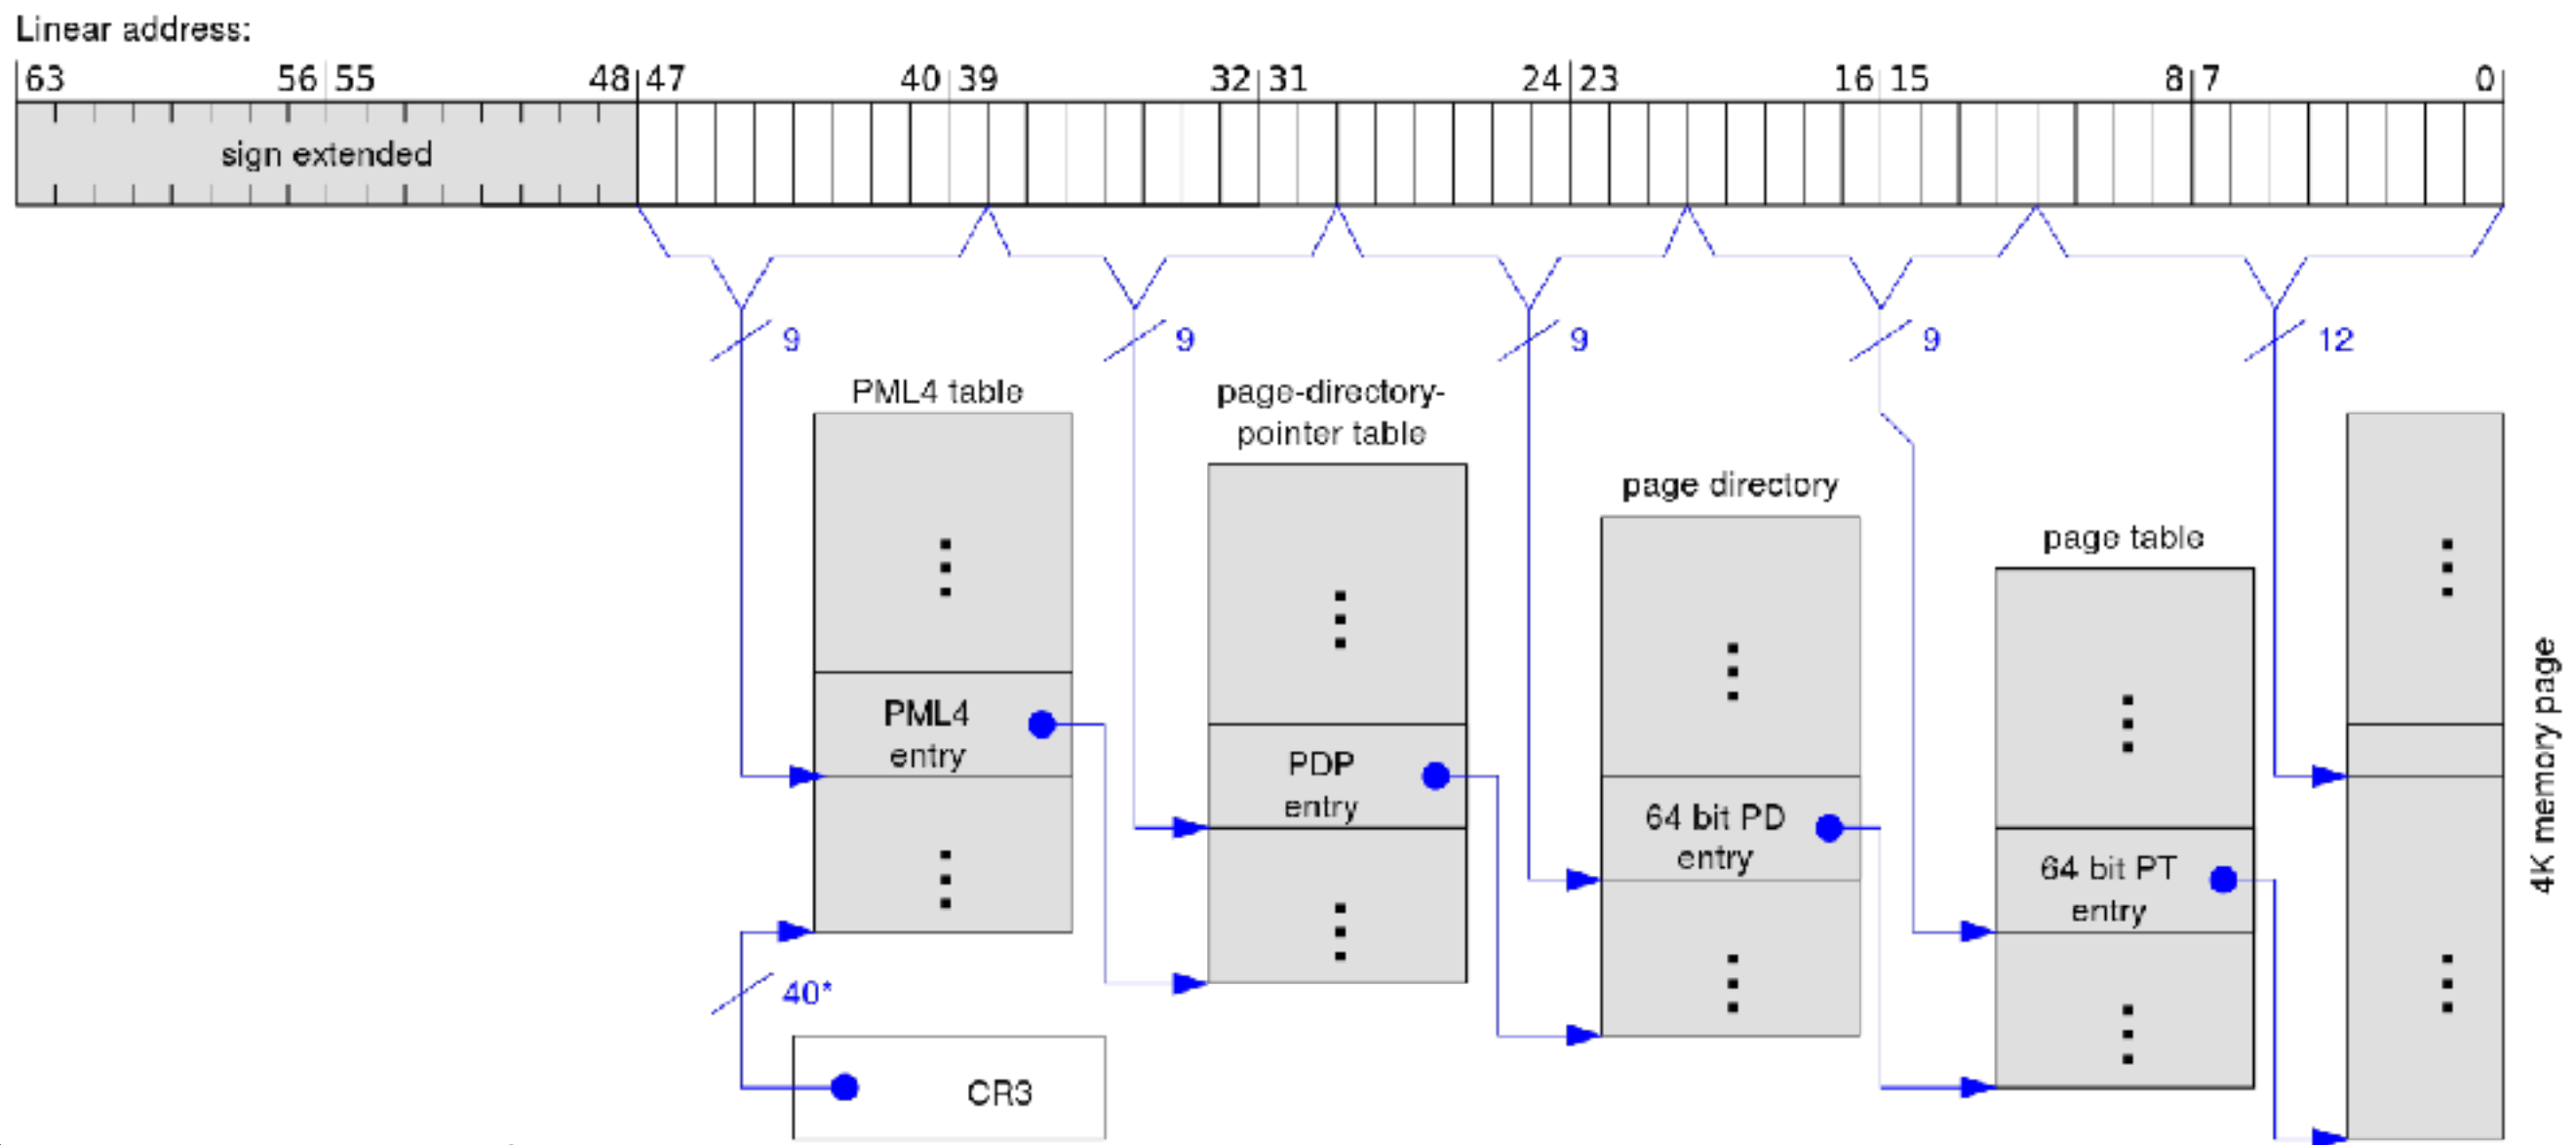
- Never heard of a program using 256TB of memory
  - Page table is a sparse data structure
  - Should be divided, with pieces allocated on demand!



\* source:Wikipedia user RokerHRO

# Multilevel Paging

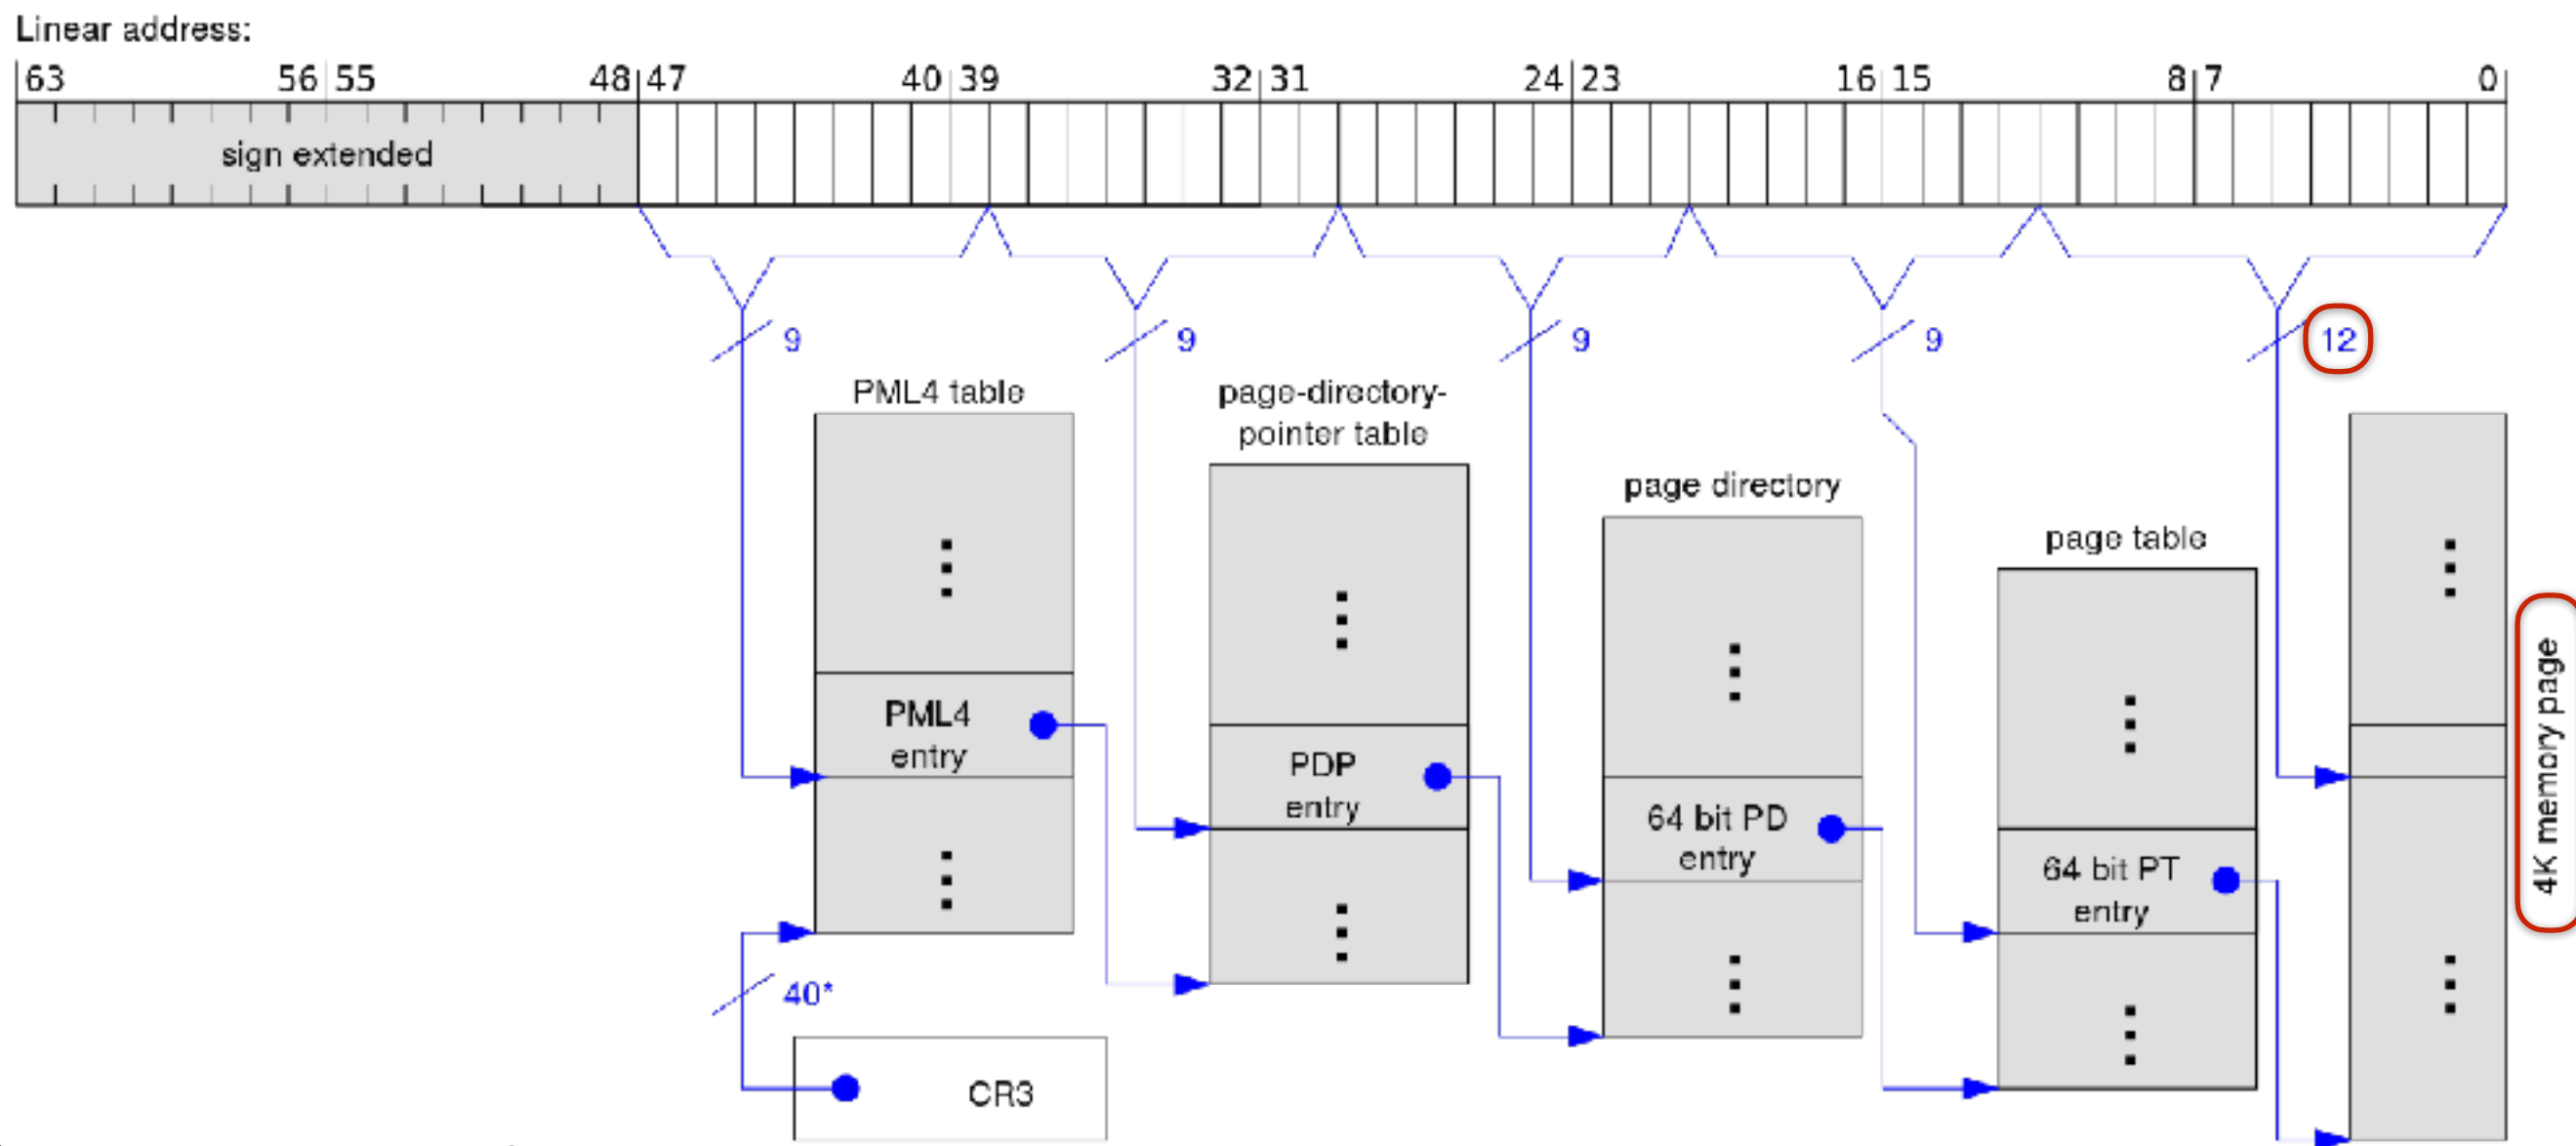
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

## Intel x86-64

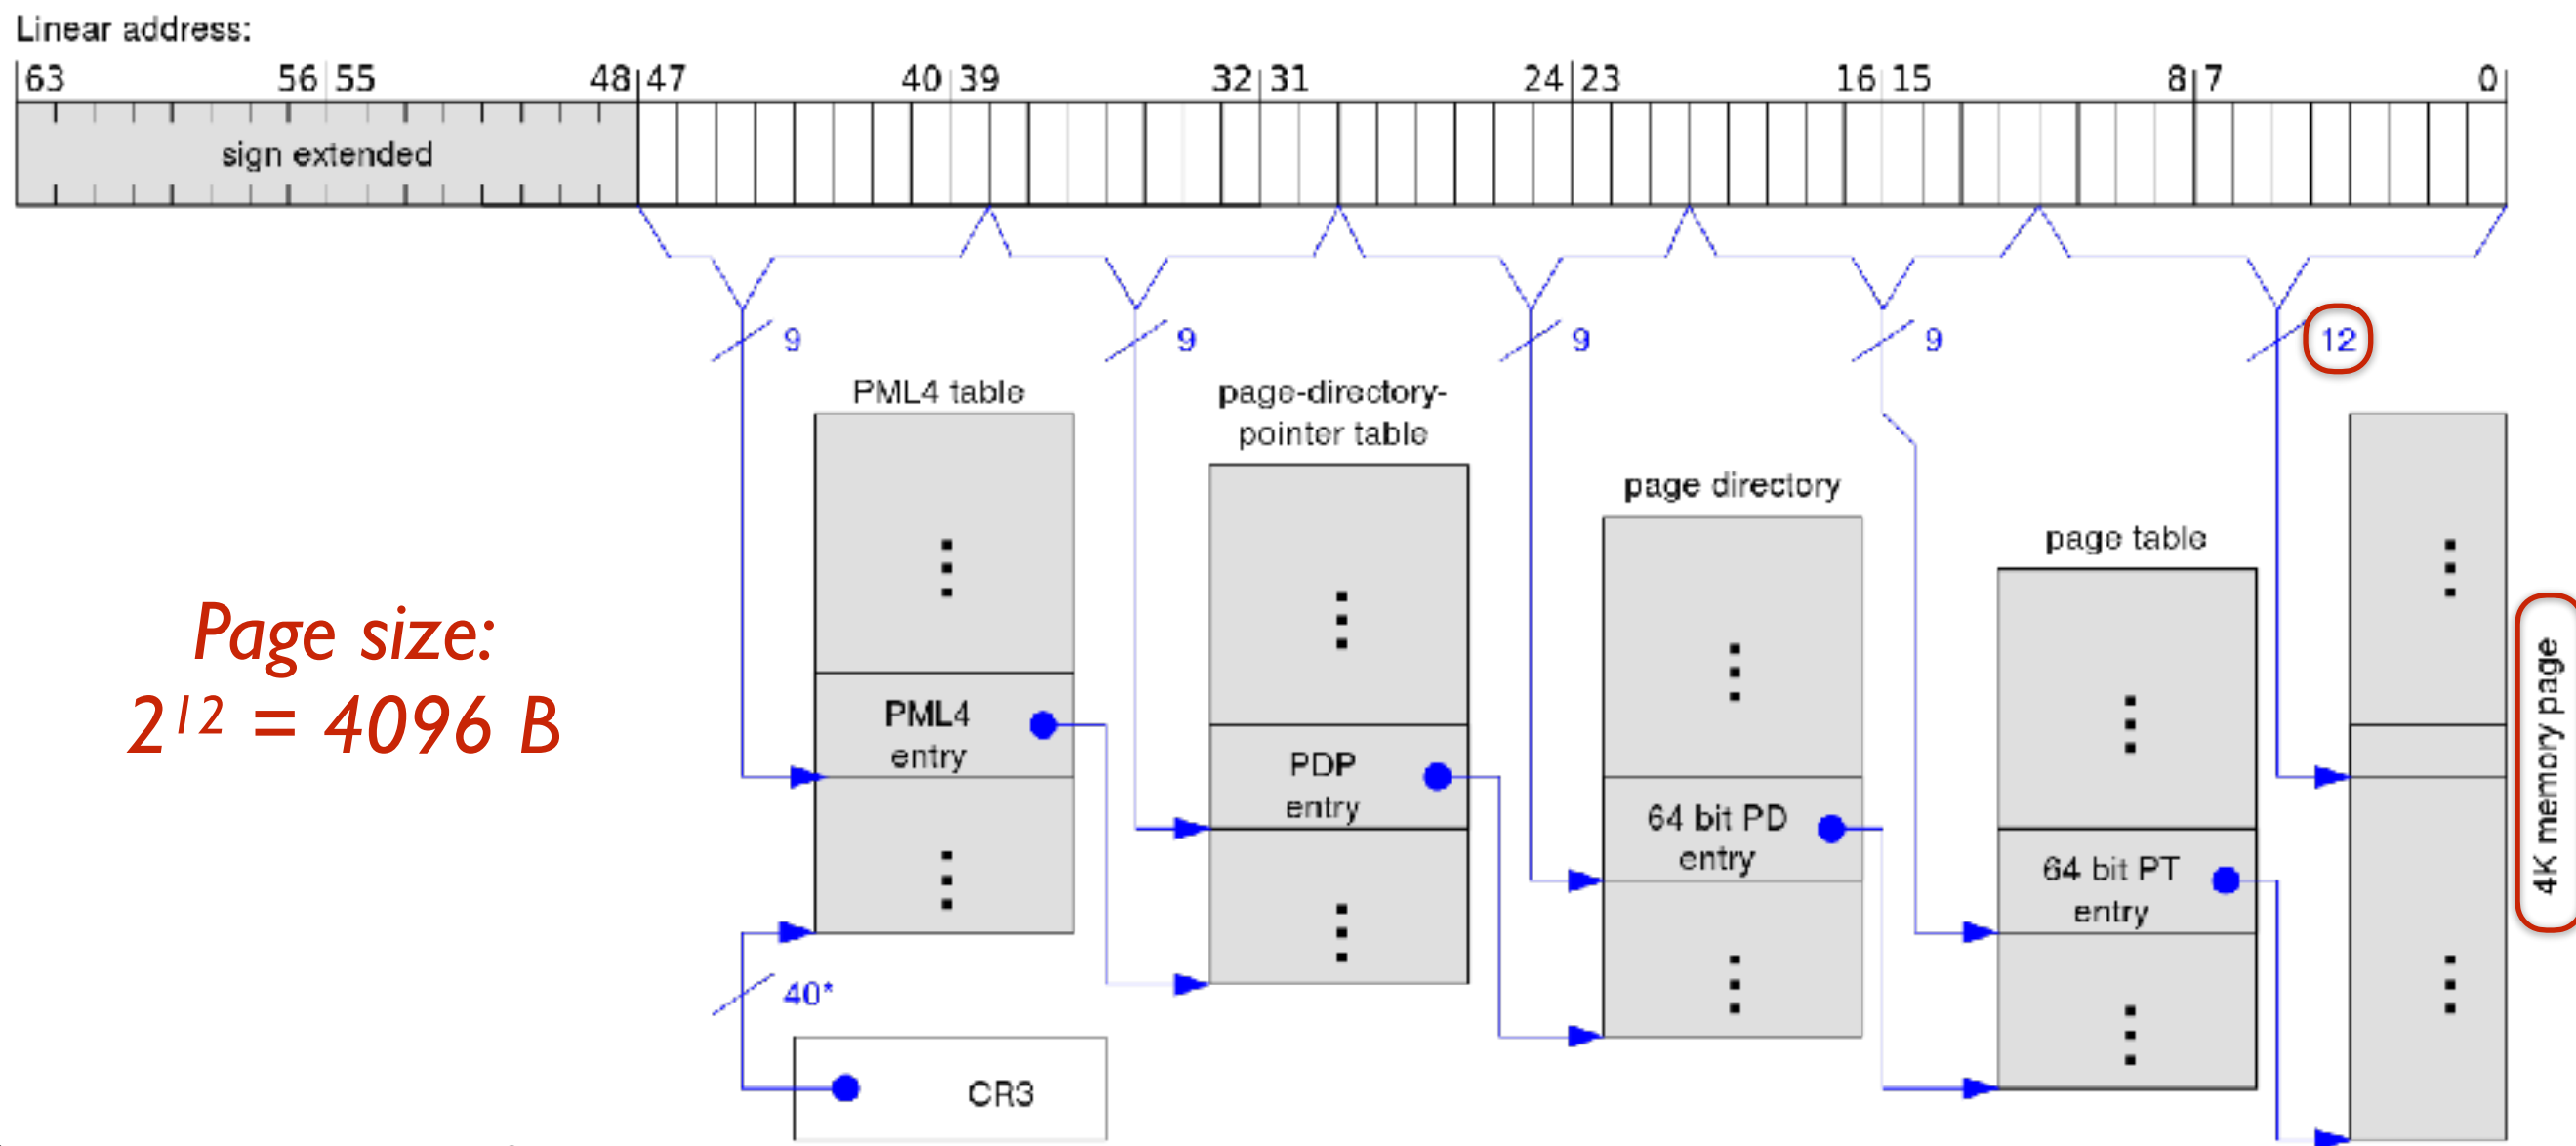


\* source: Wikipedia user RokerHRO



# Multilevel Paging

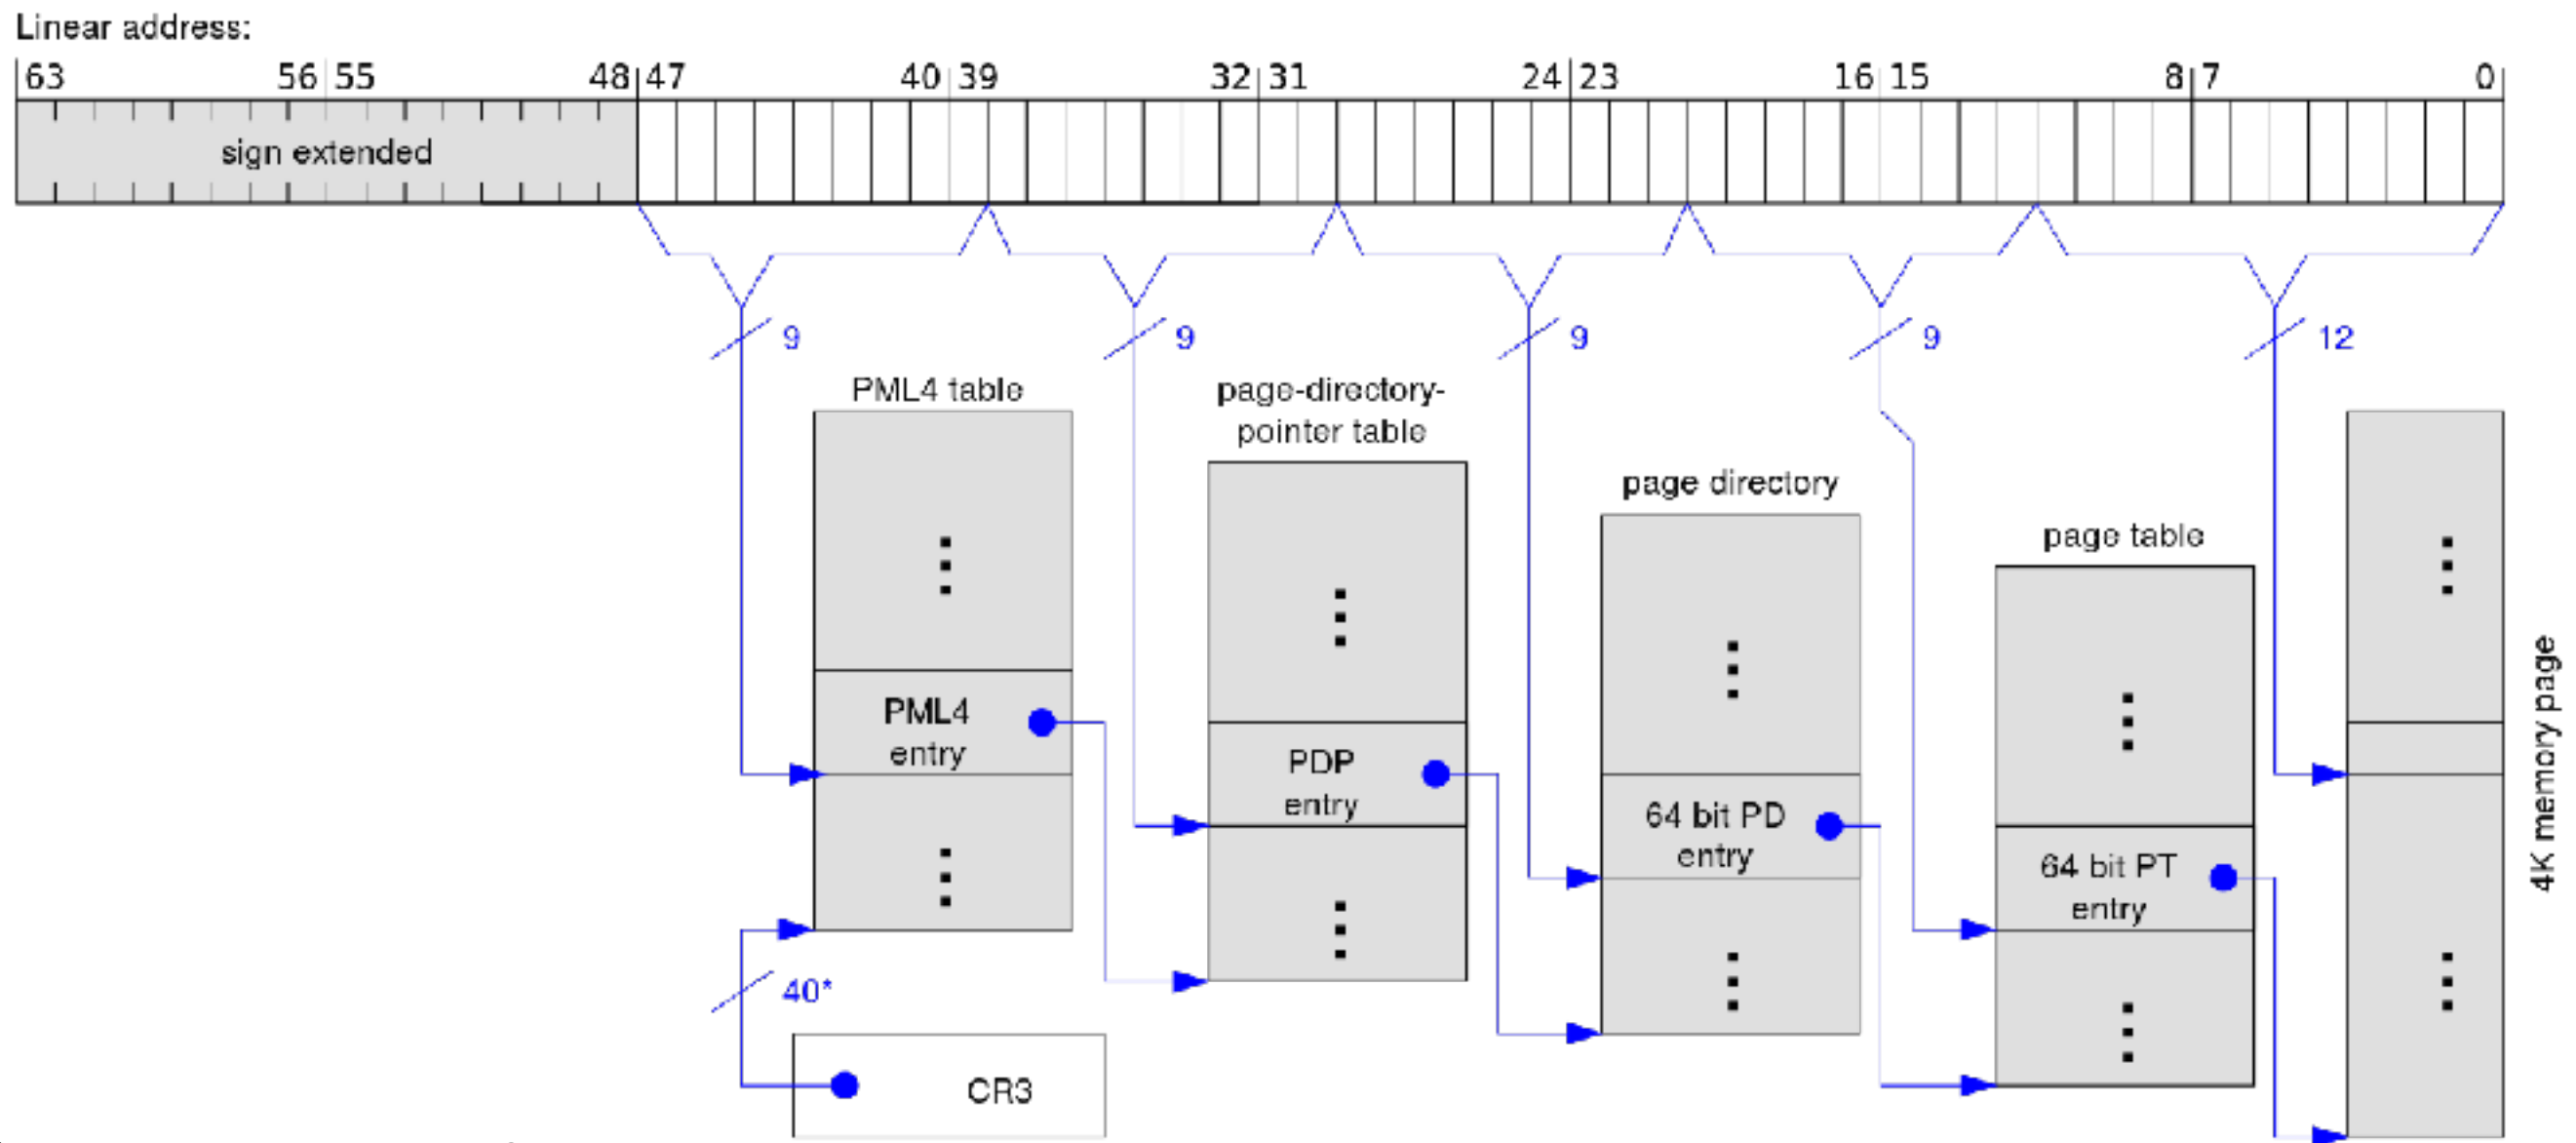
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

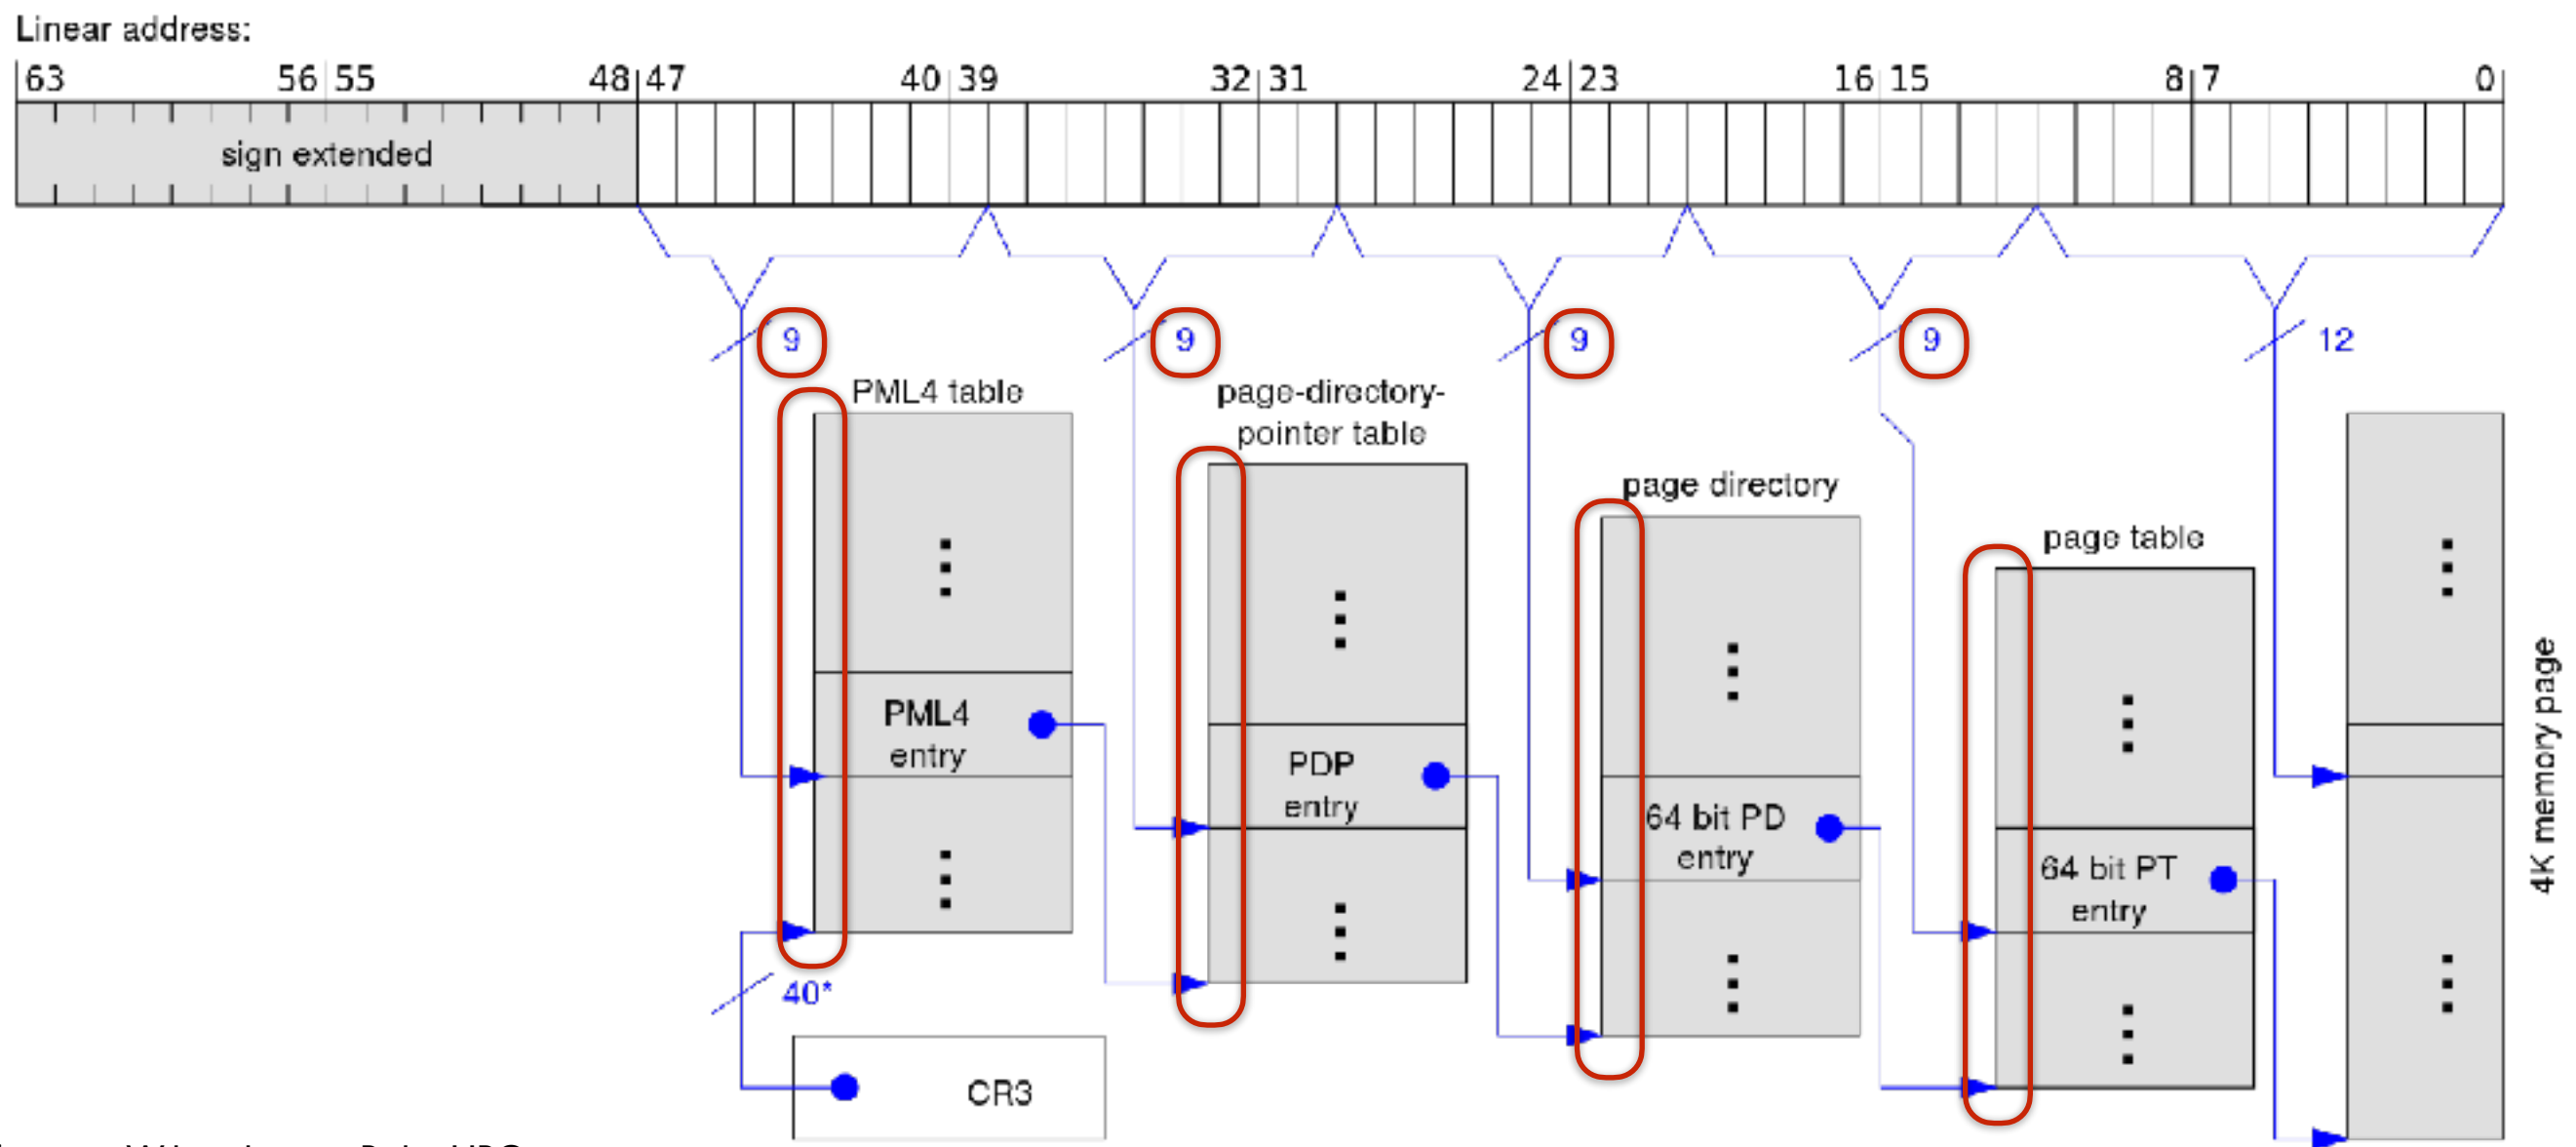
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

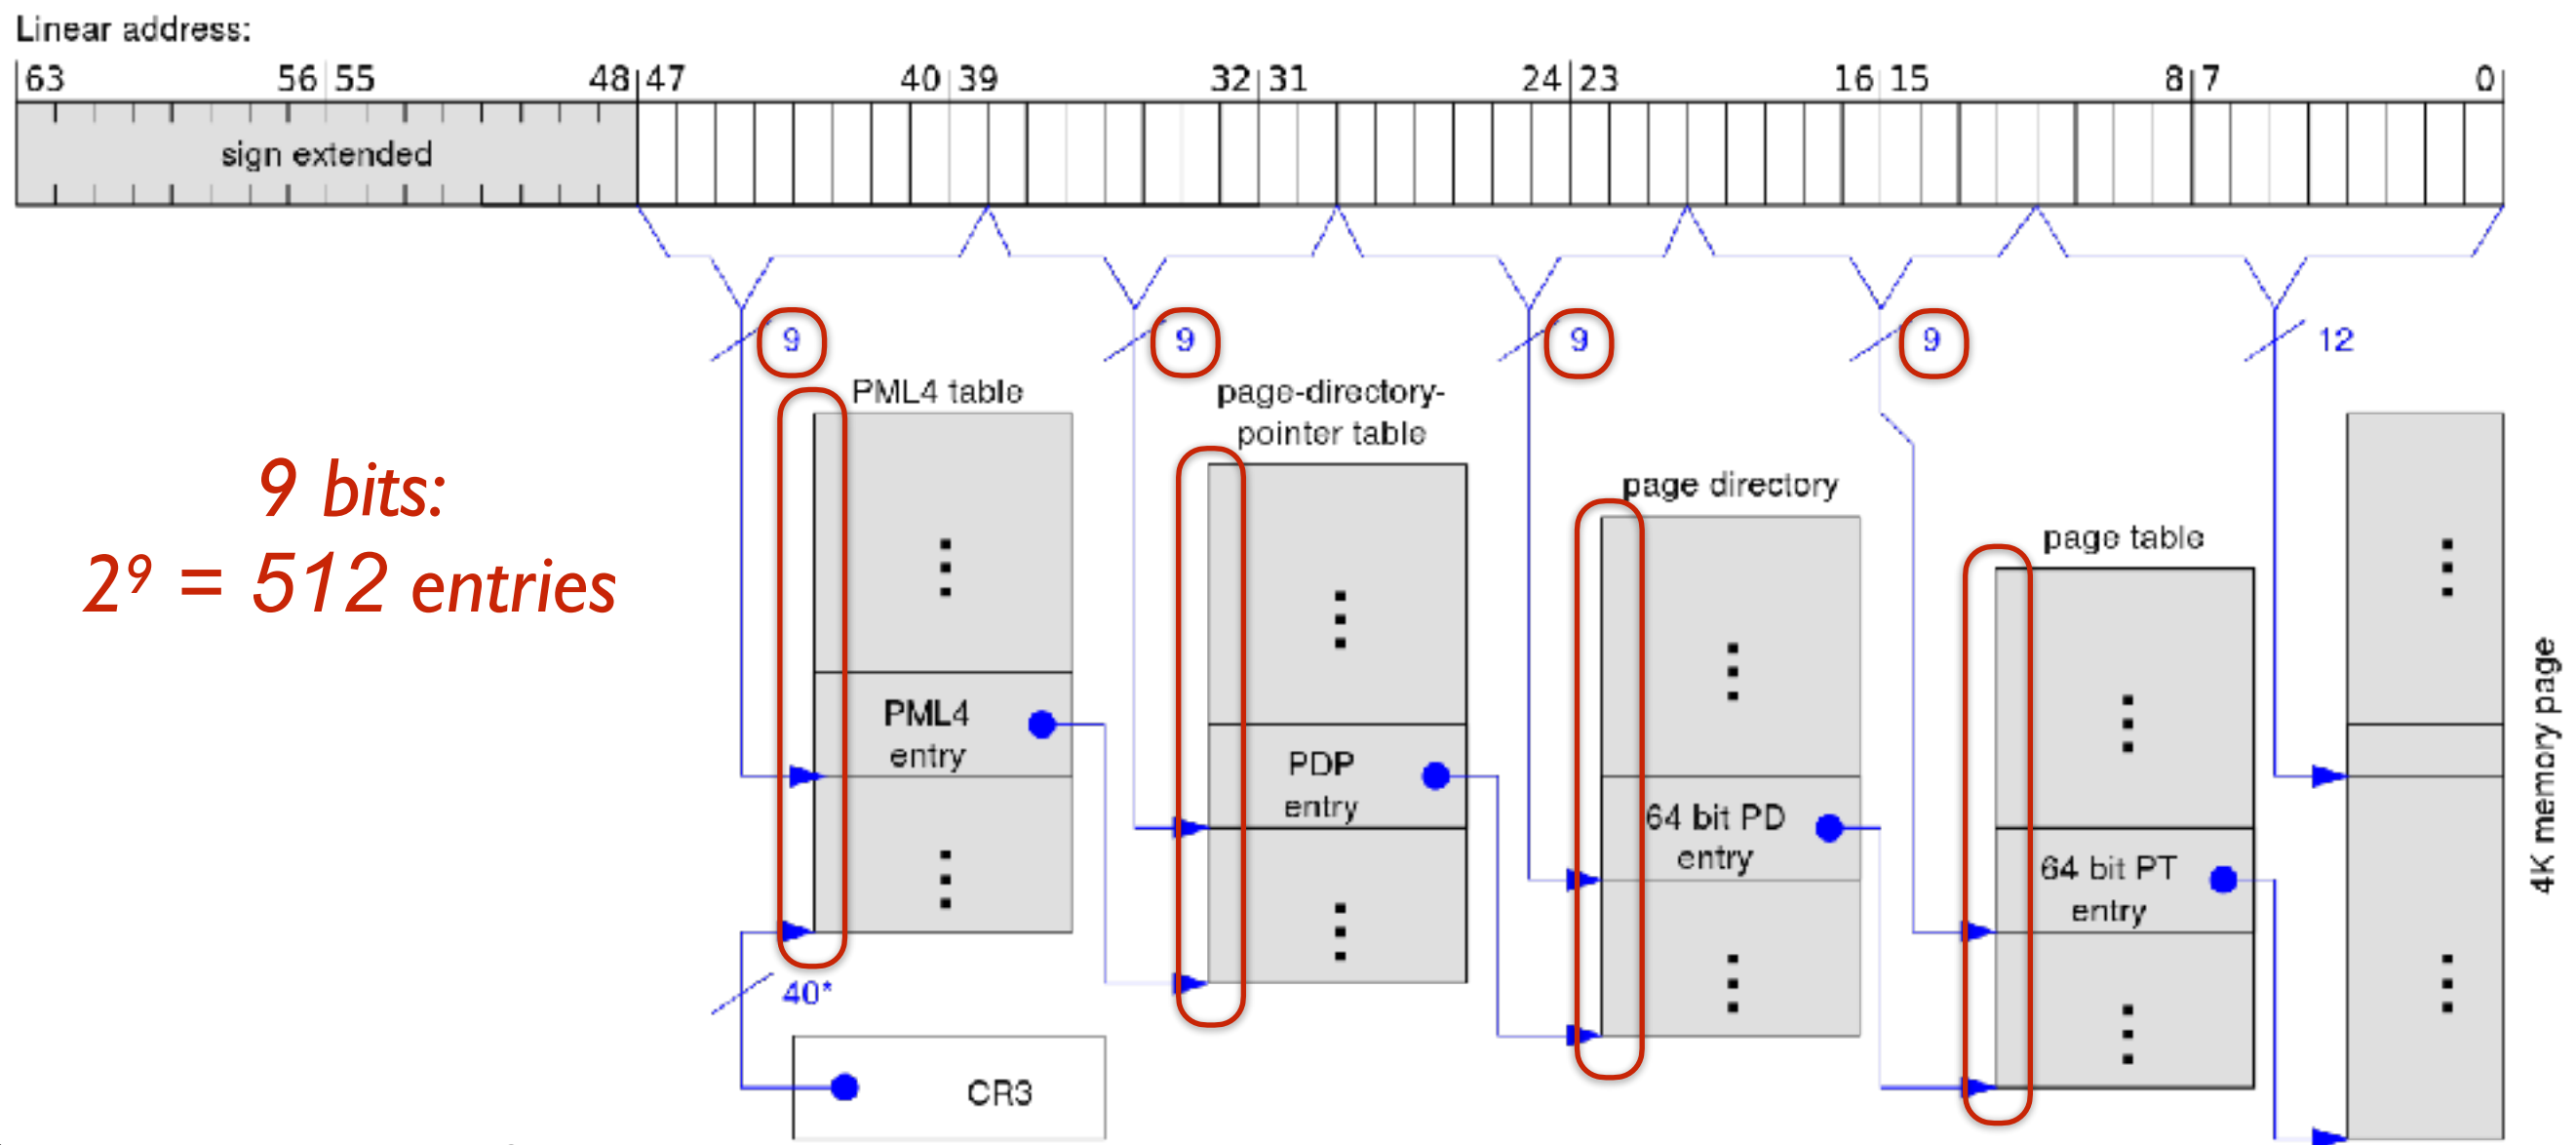
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

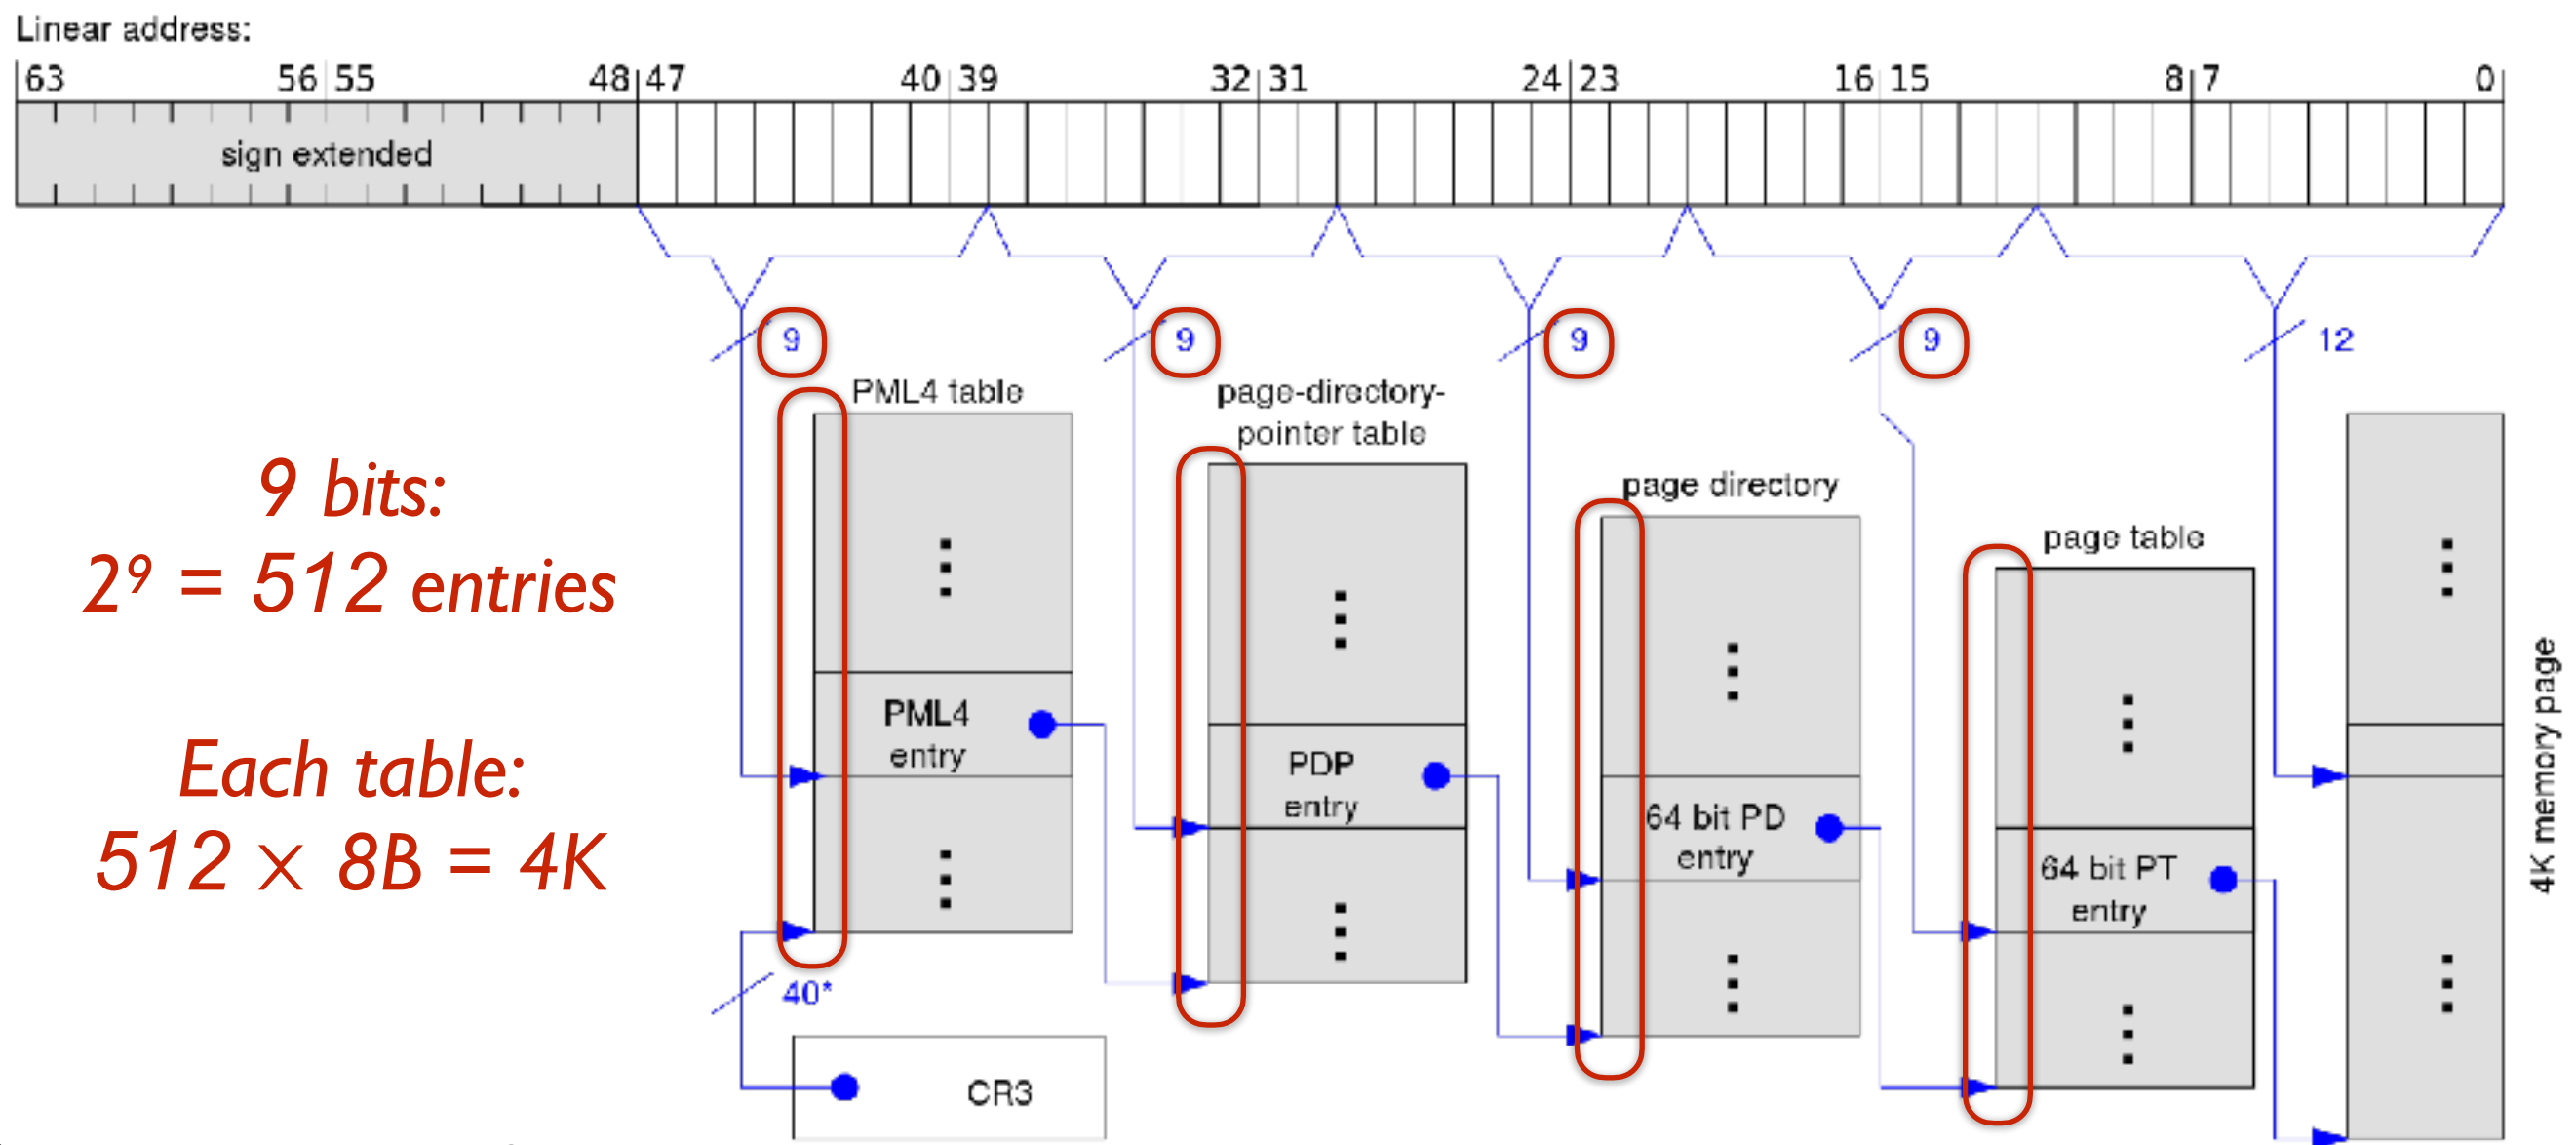
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging


## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

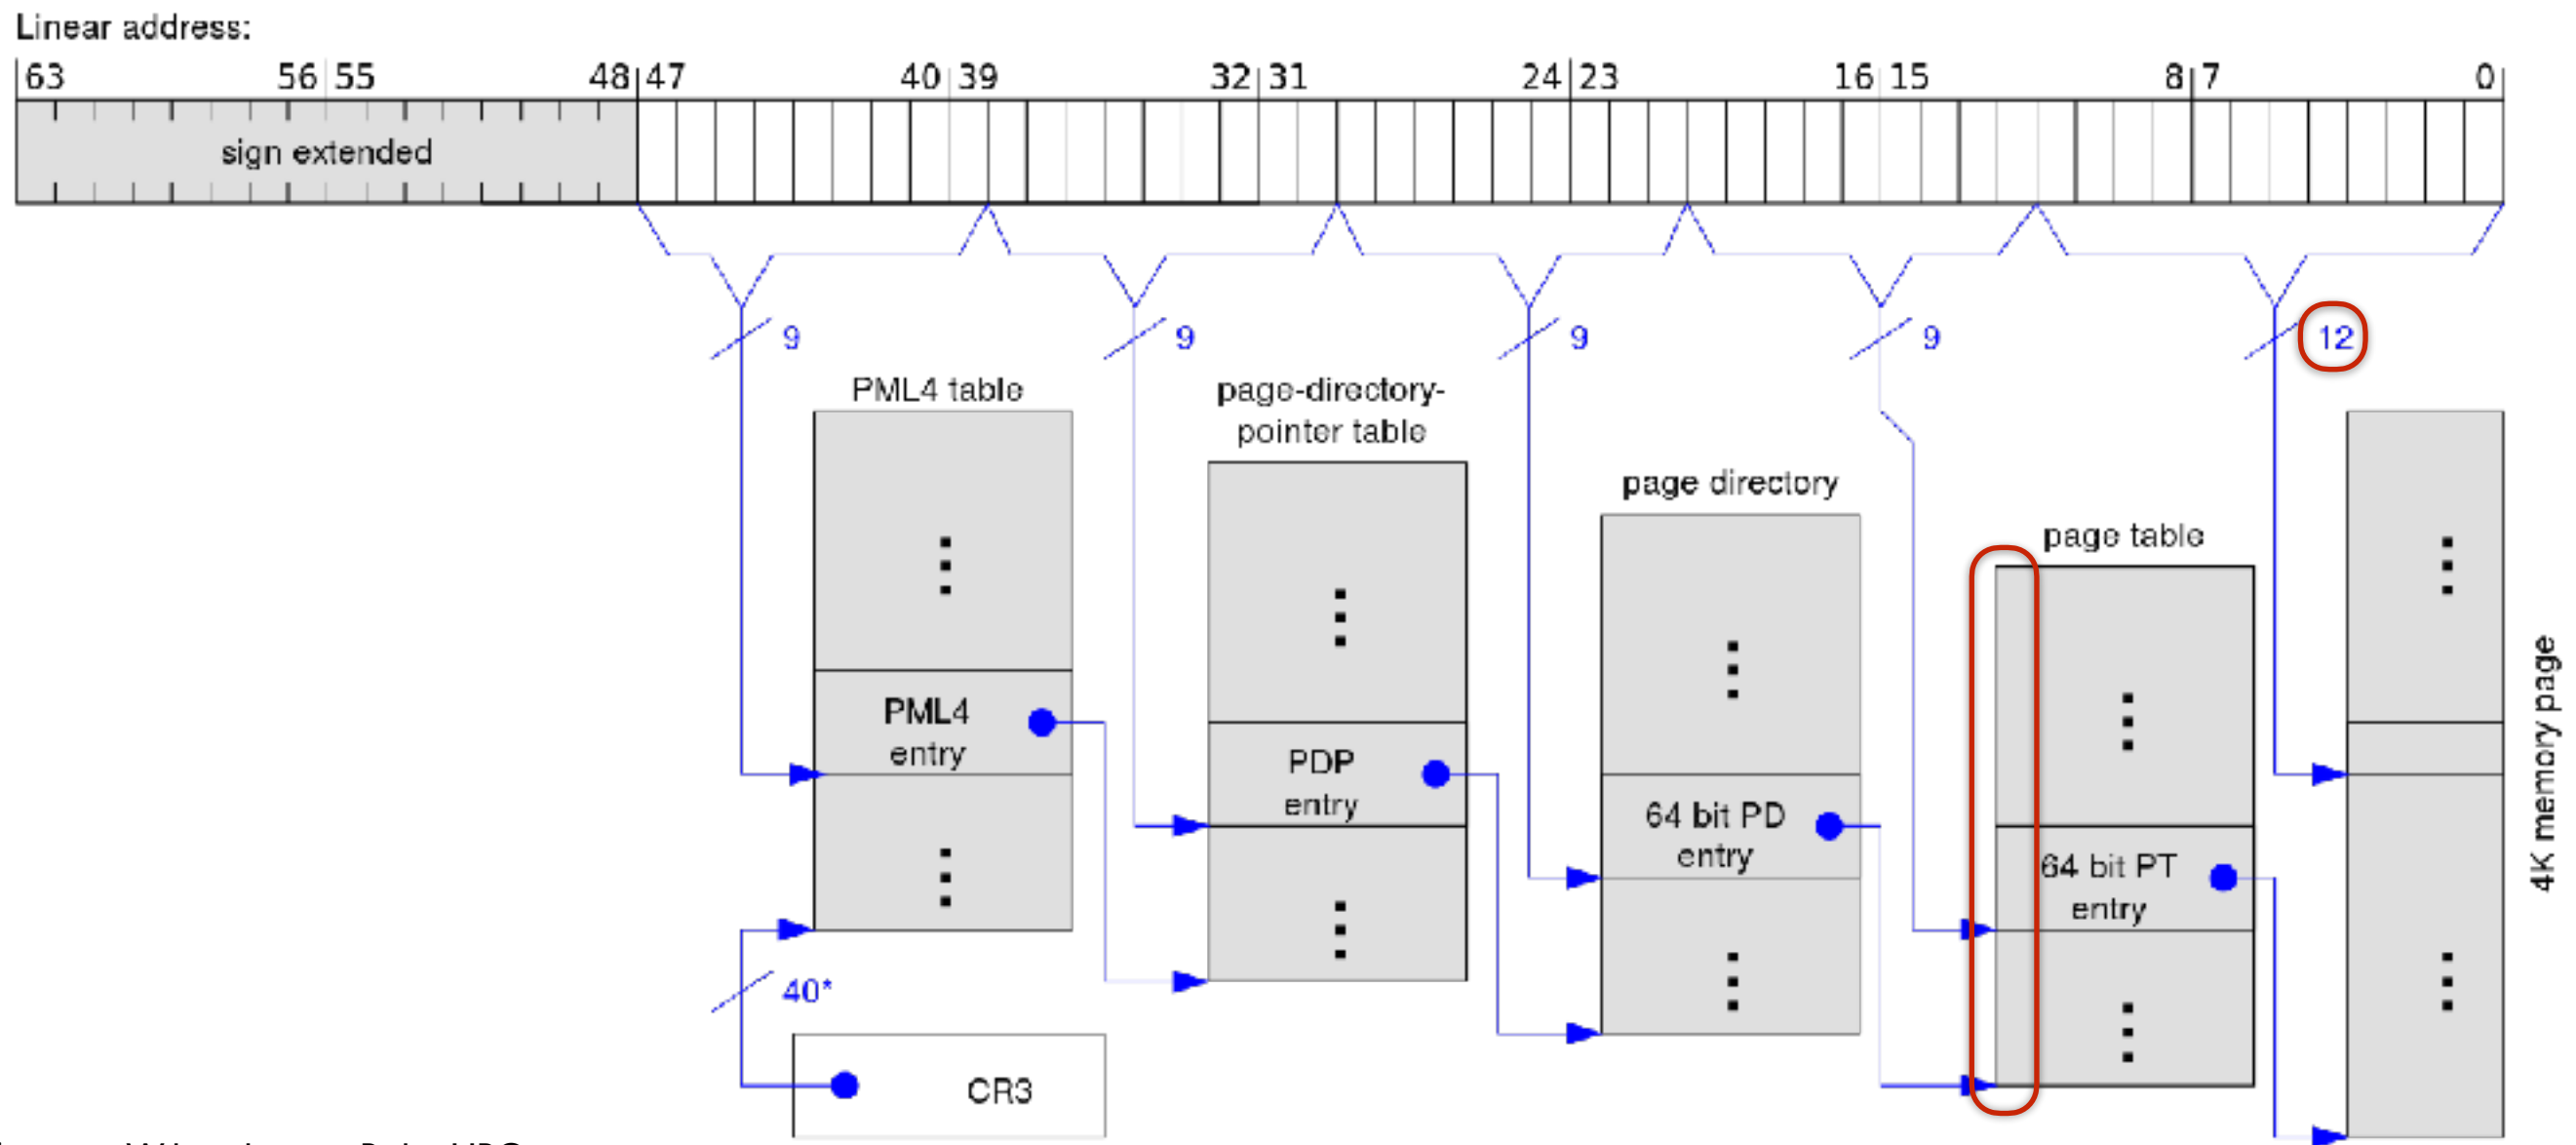
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

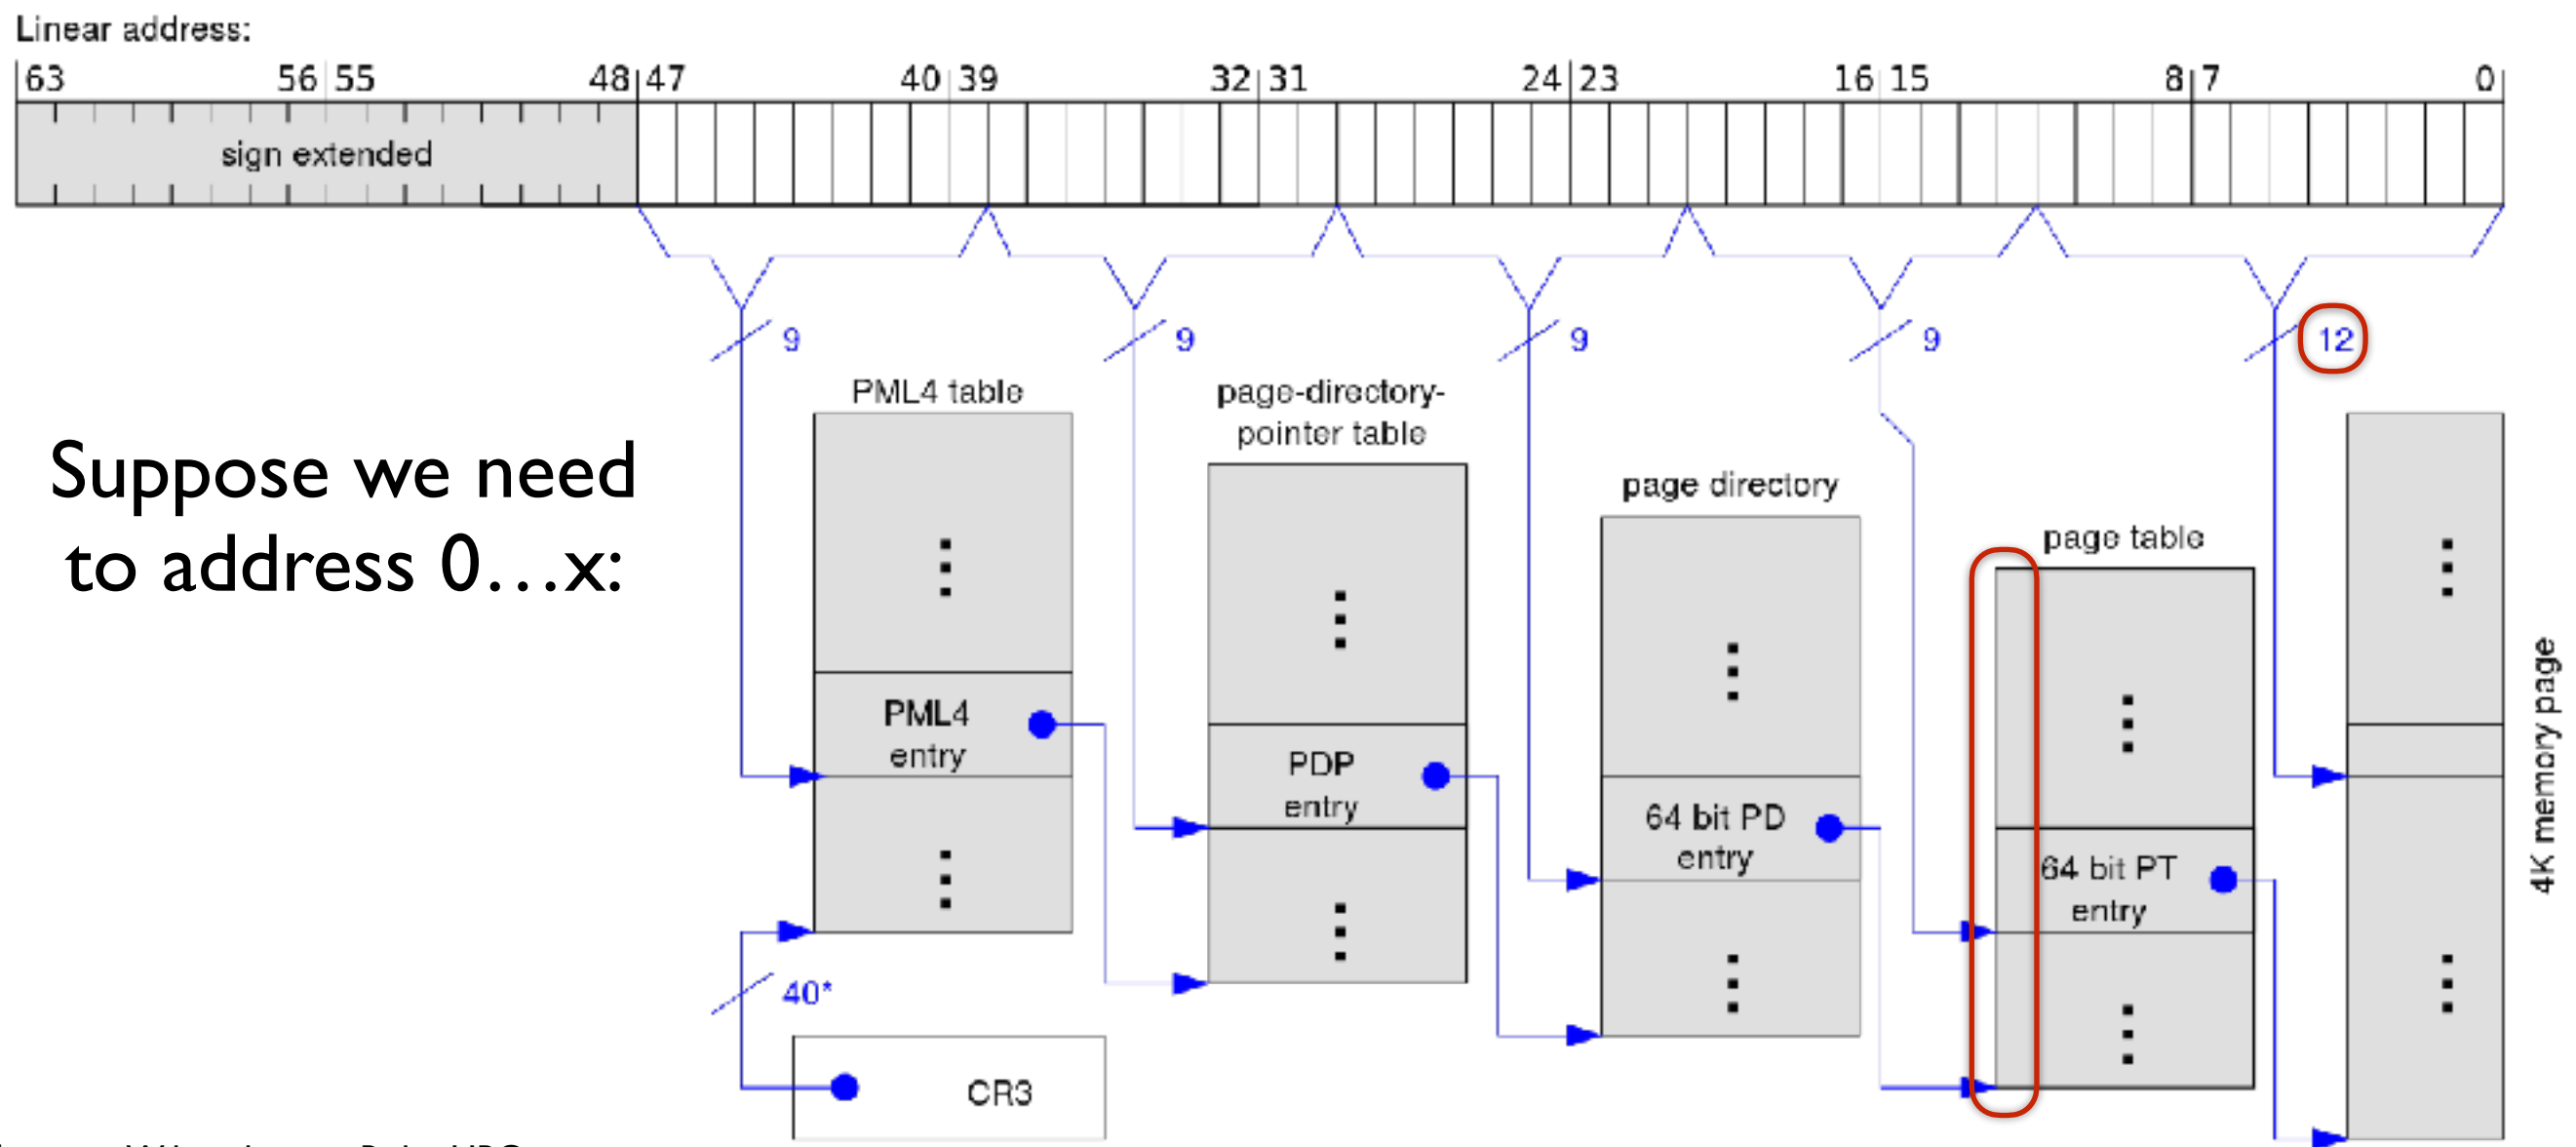
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

## Intel x86-64

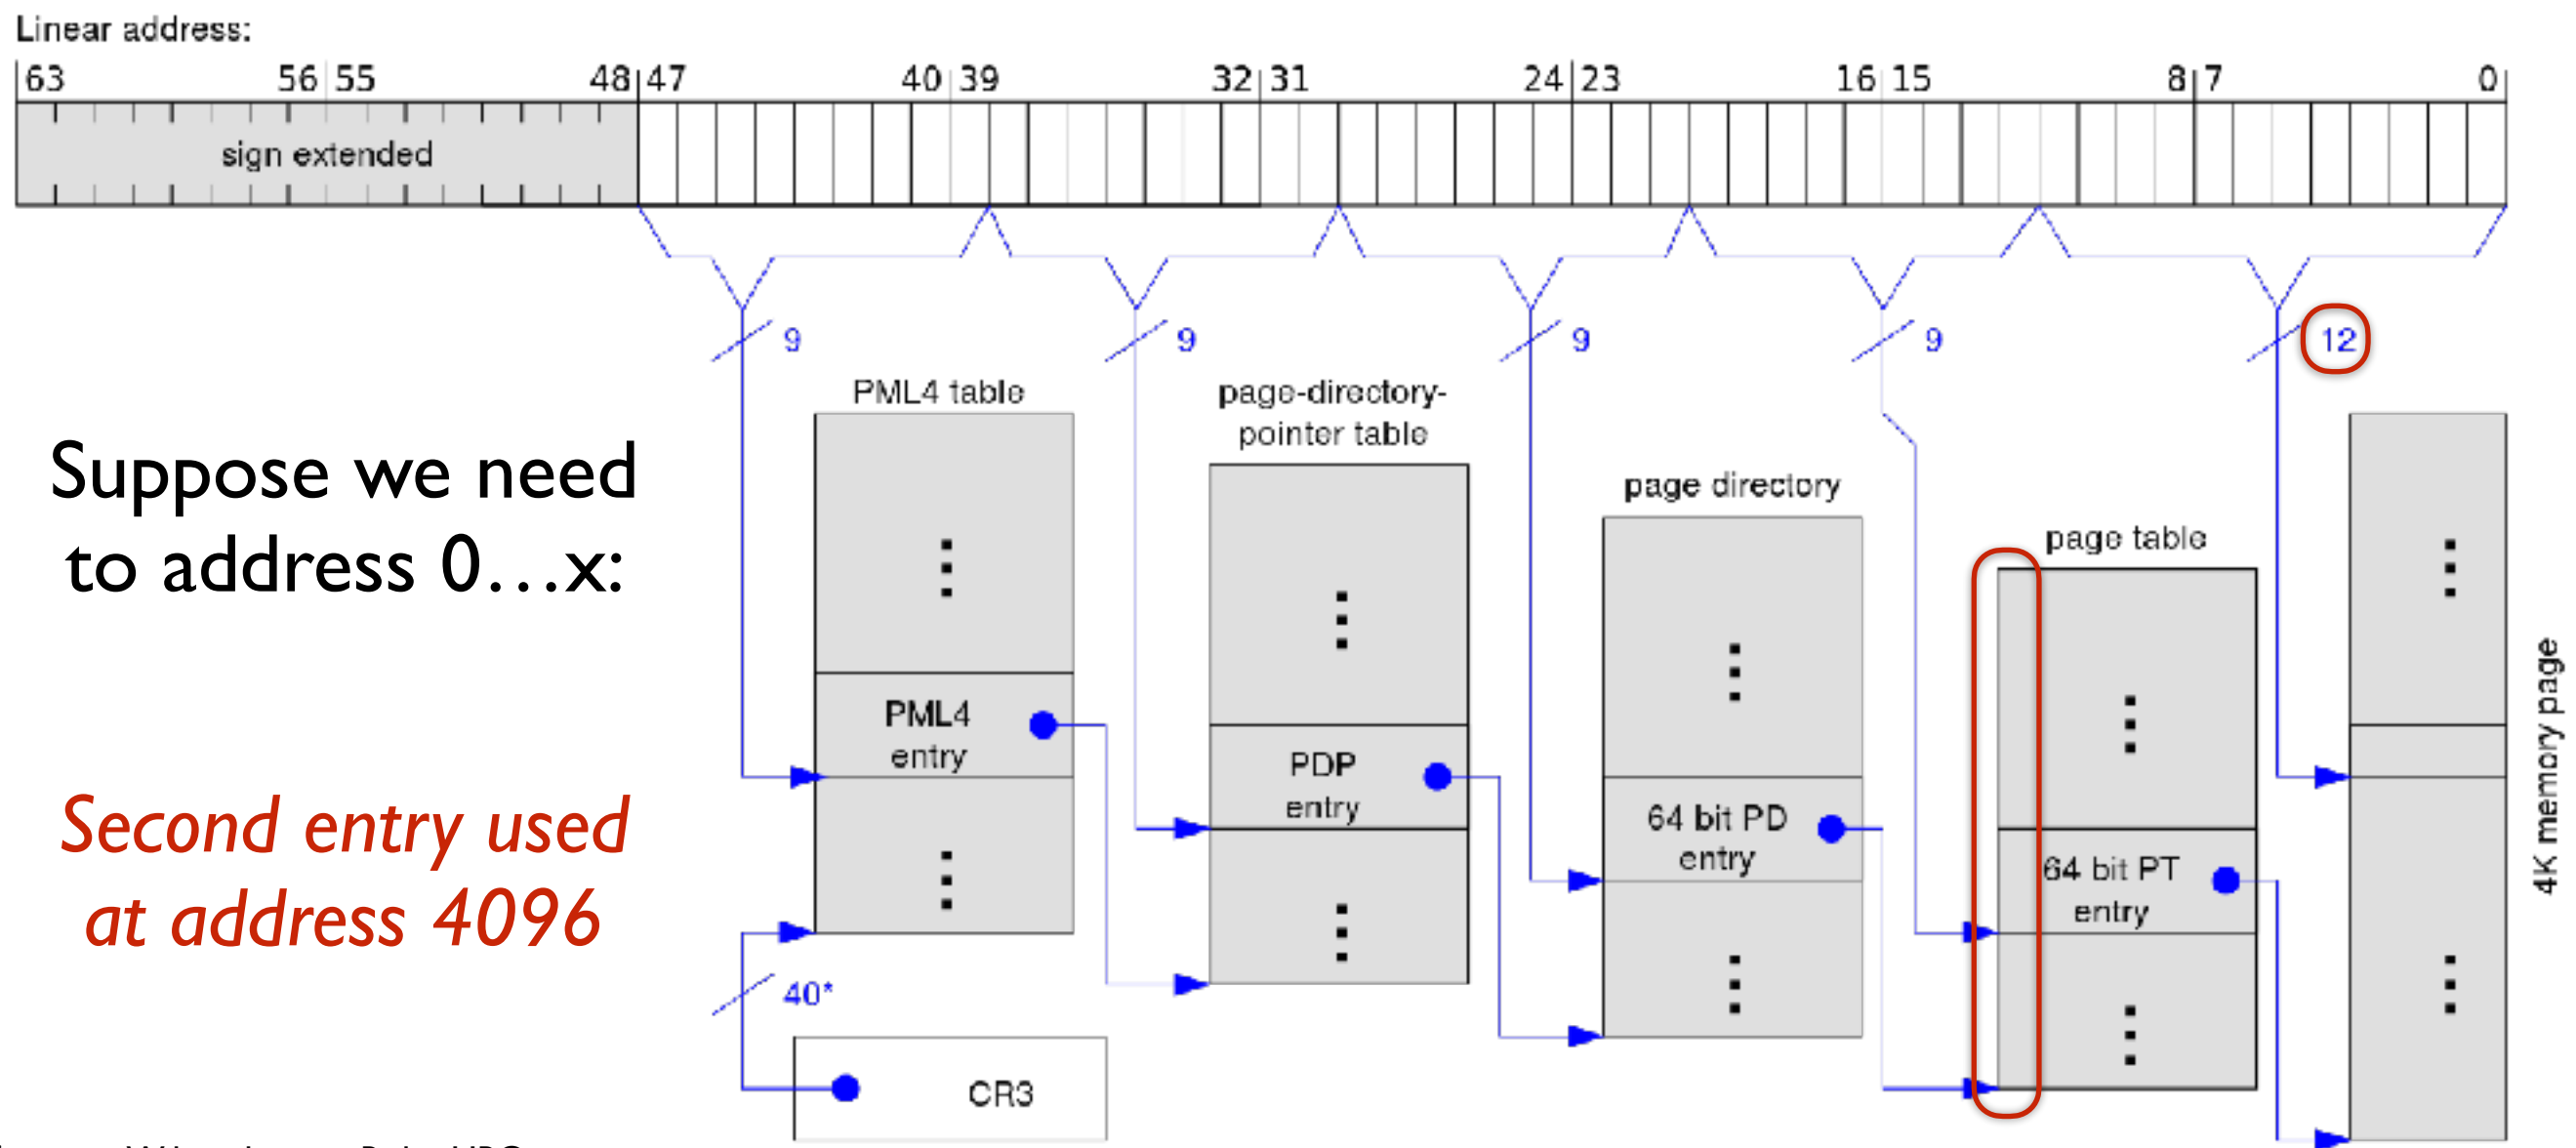


\* source: Wikipedia user RokerHRO



# Multilevel Paging

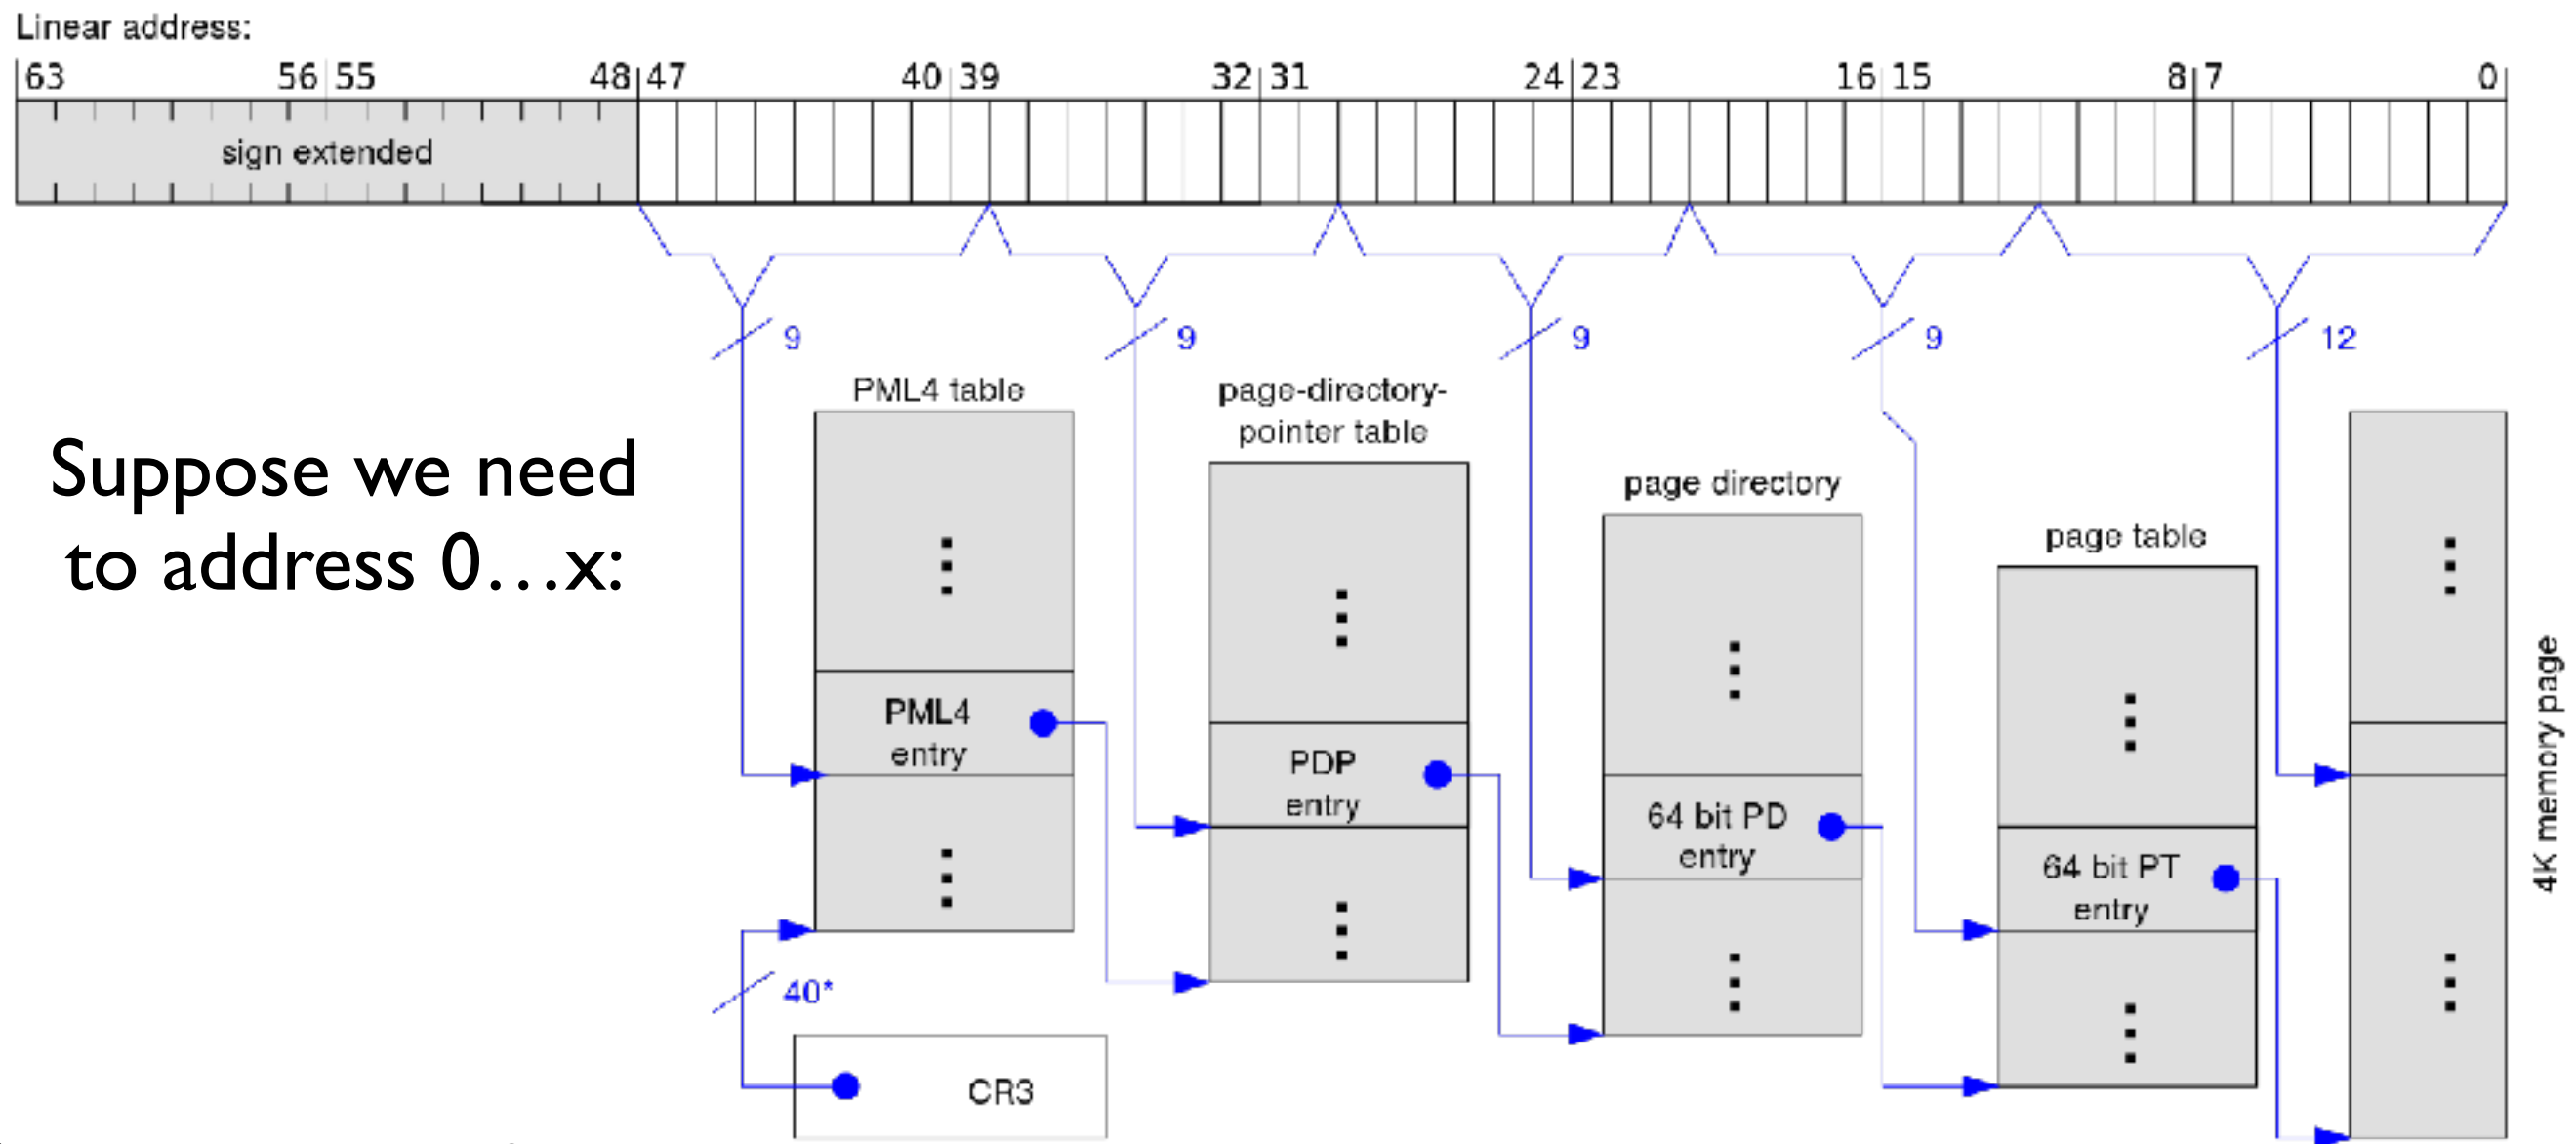
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

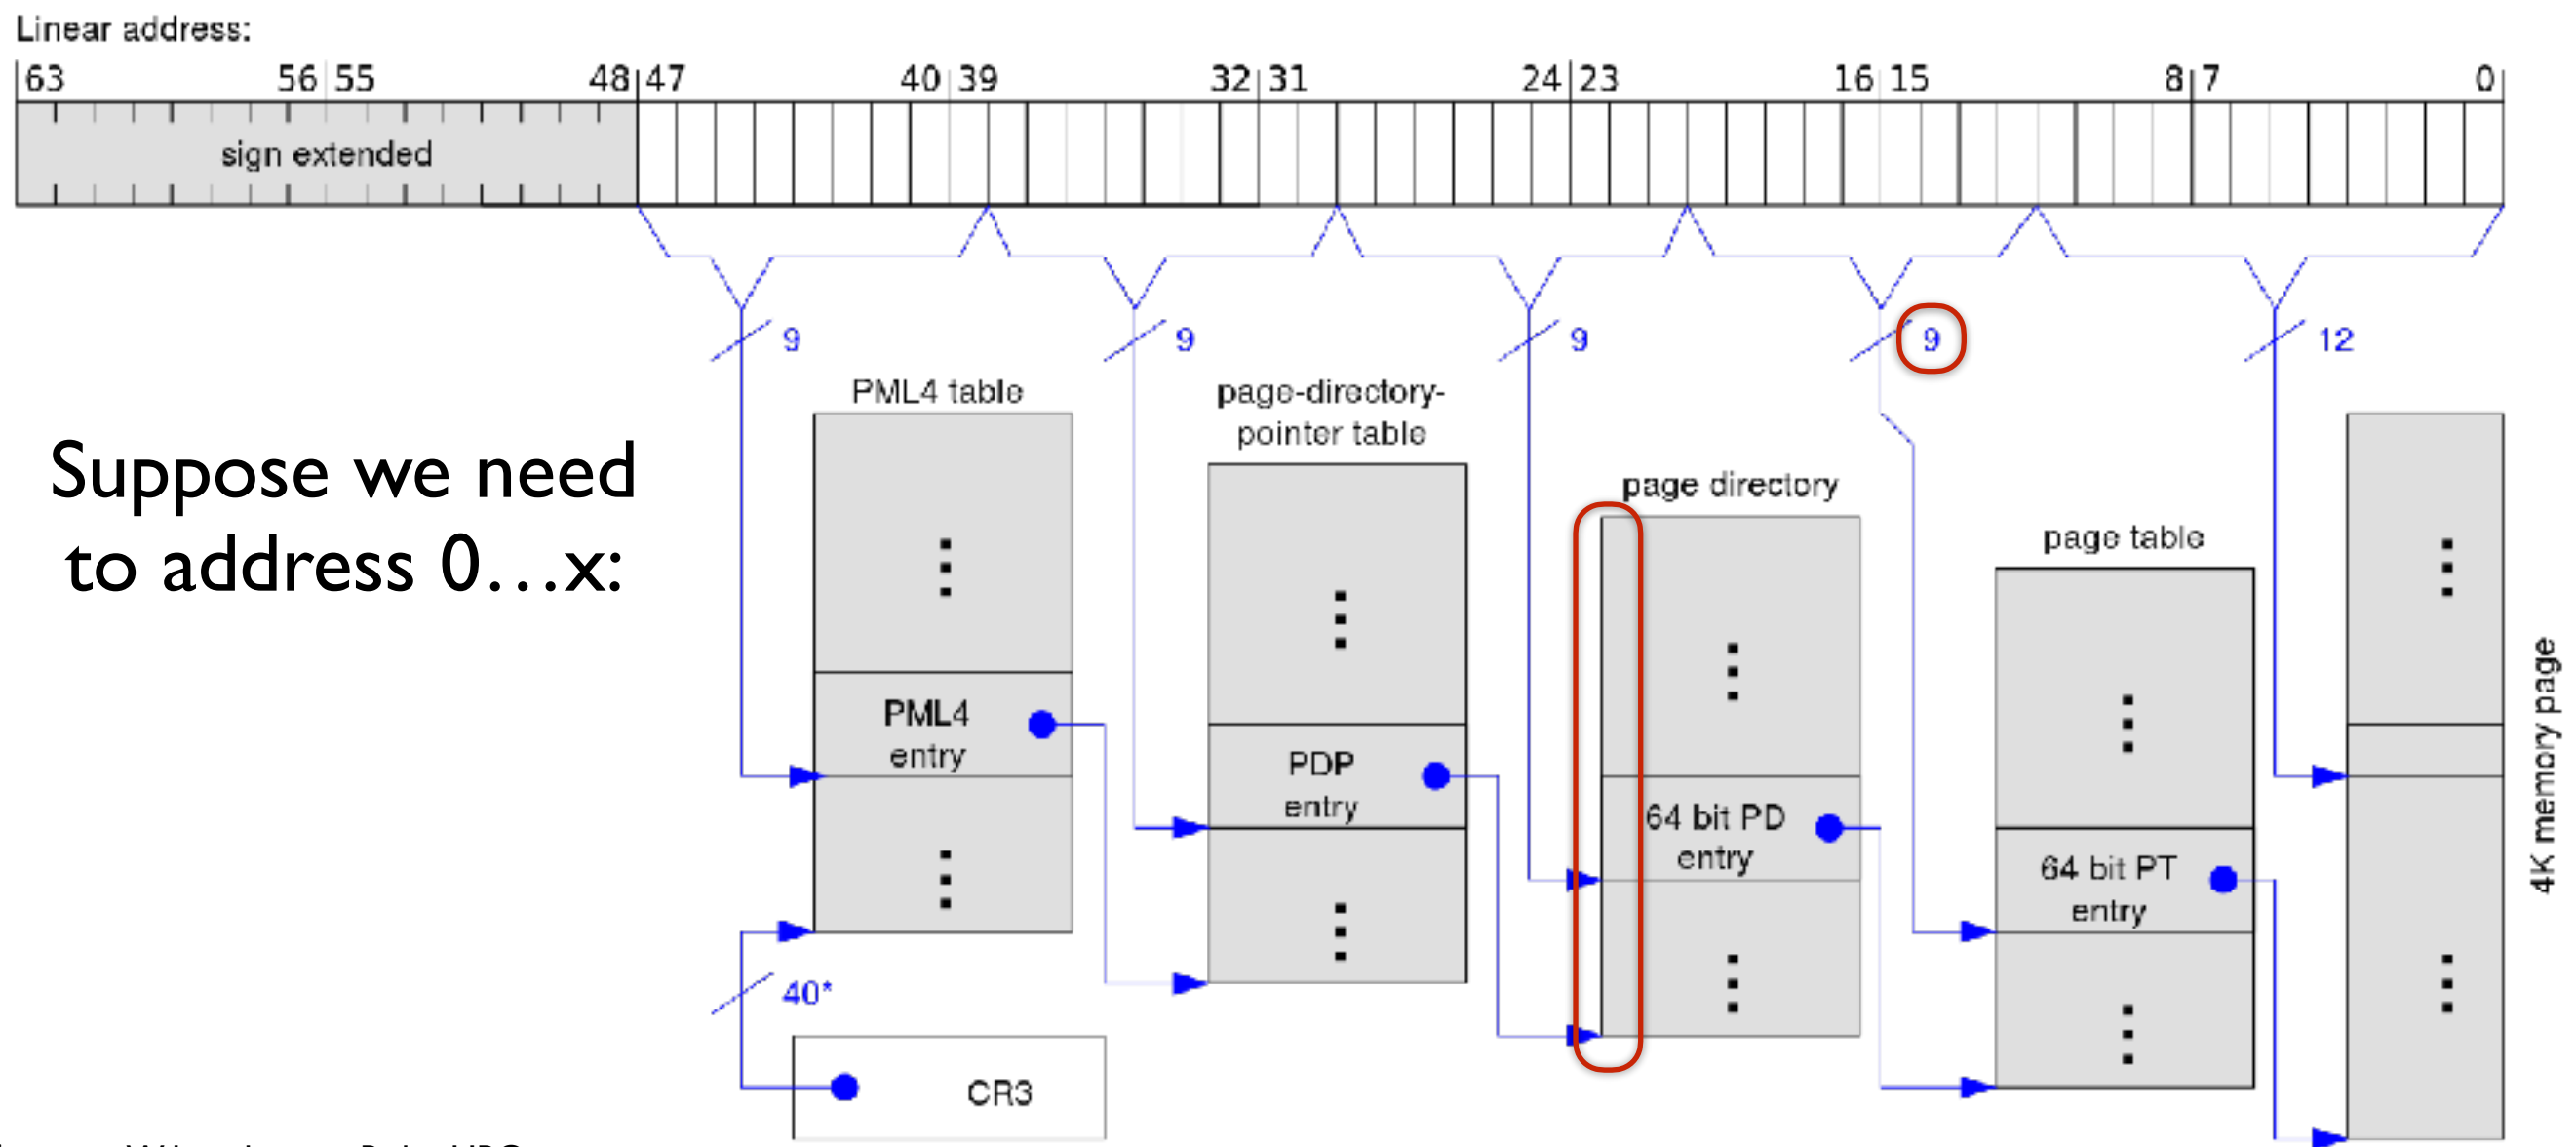
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

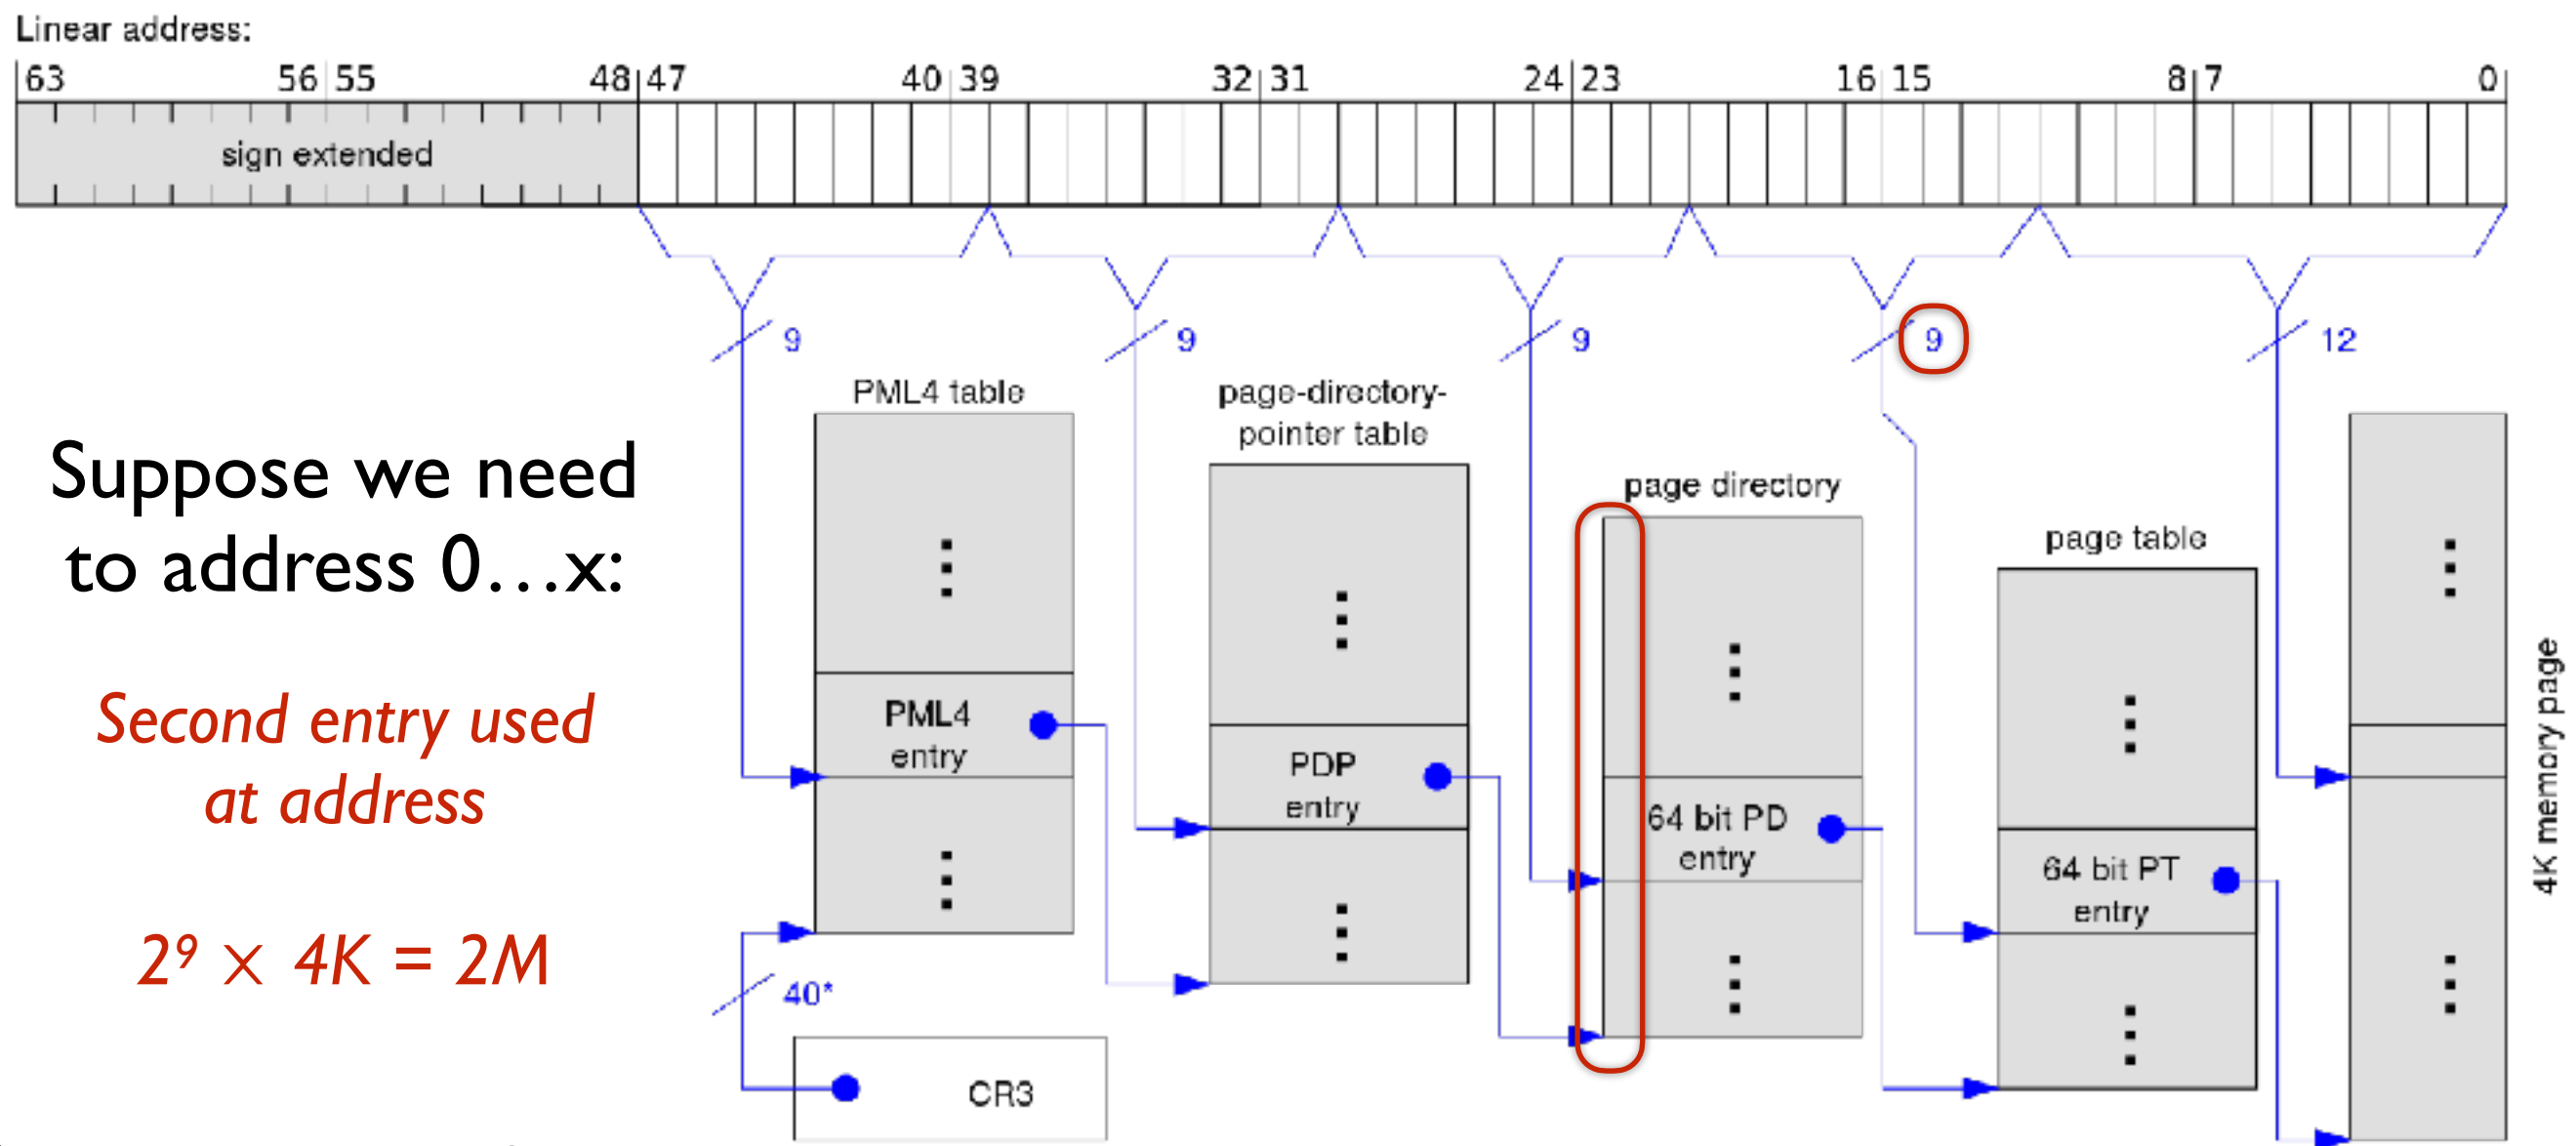
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

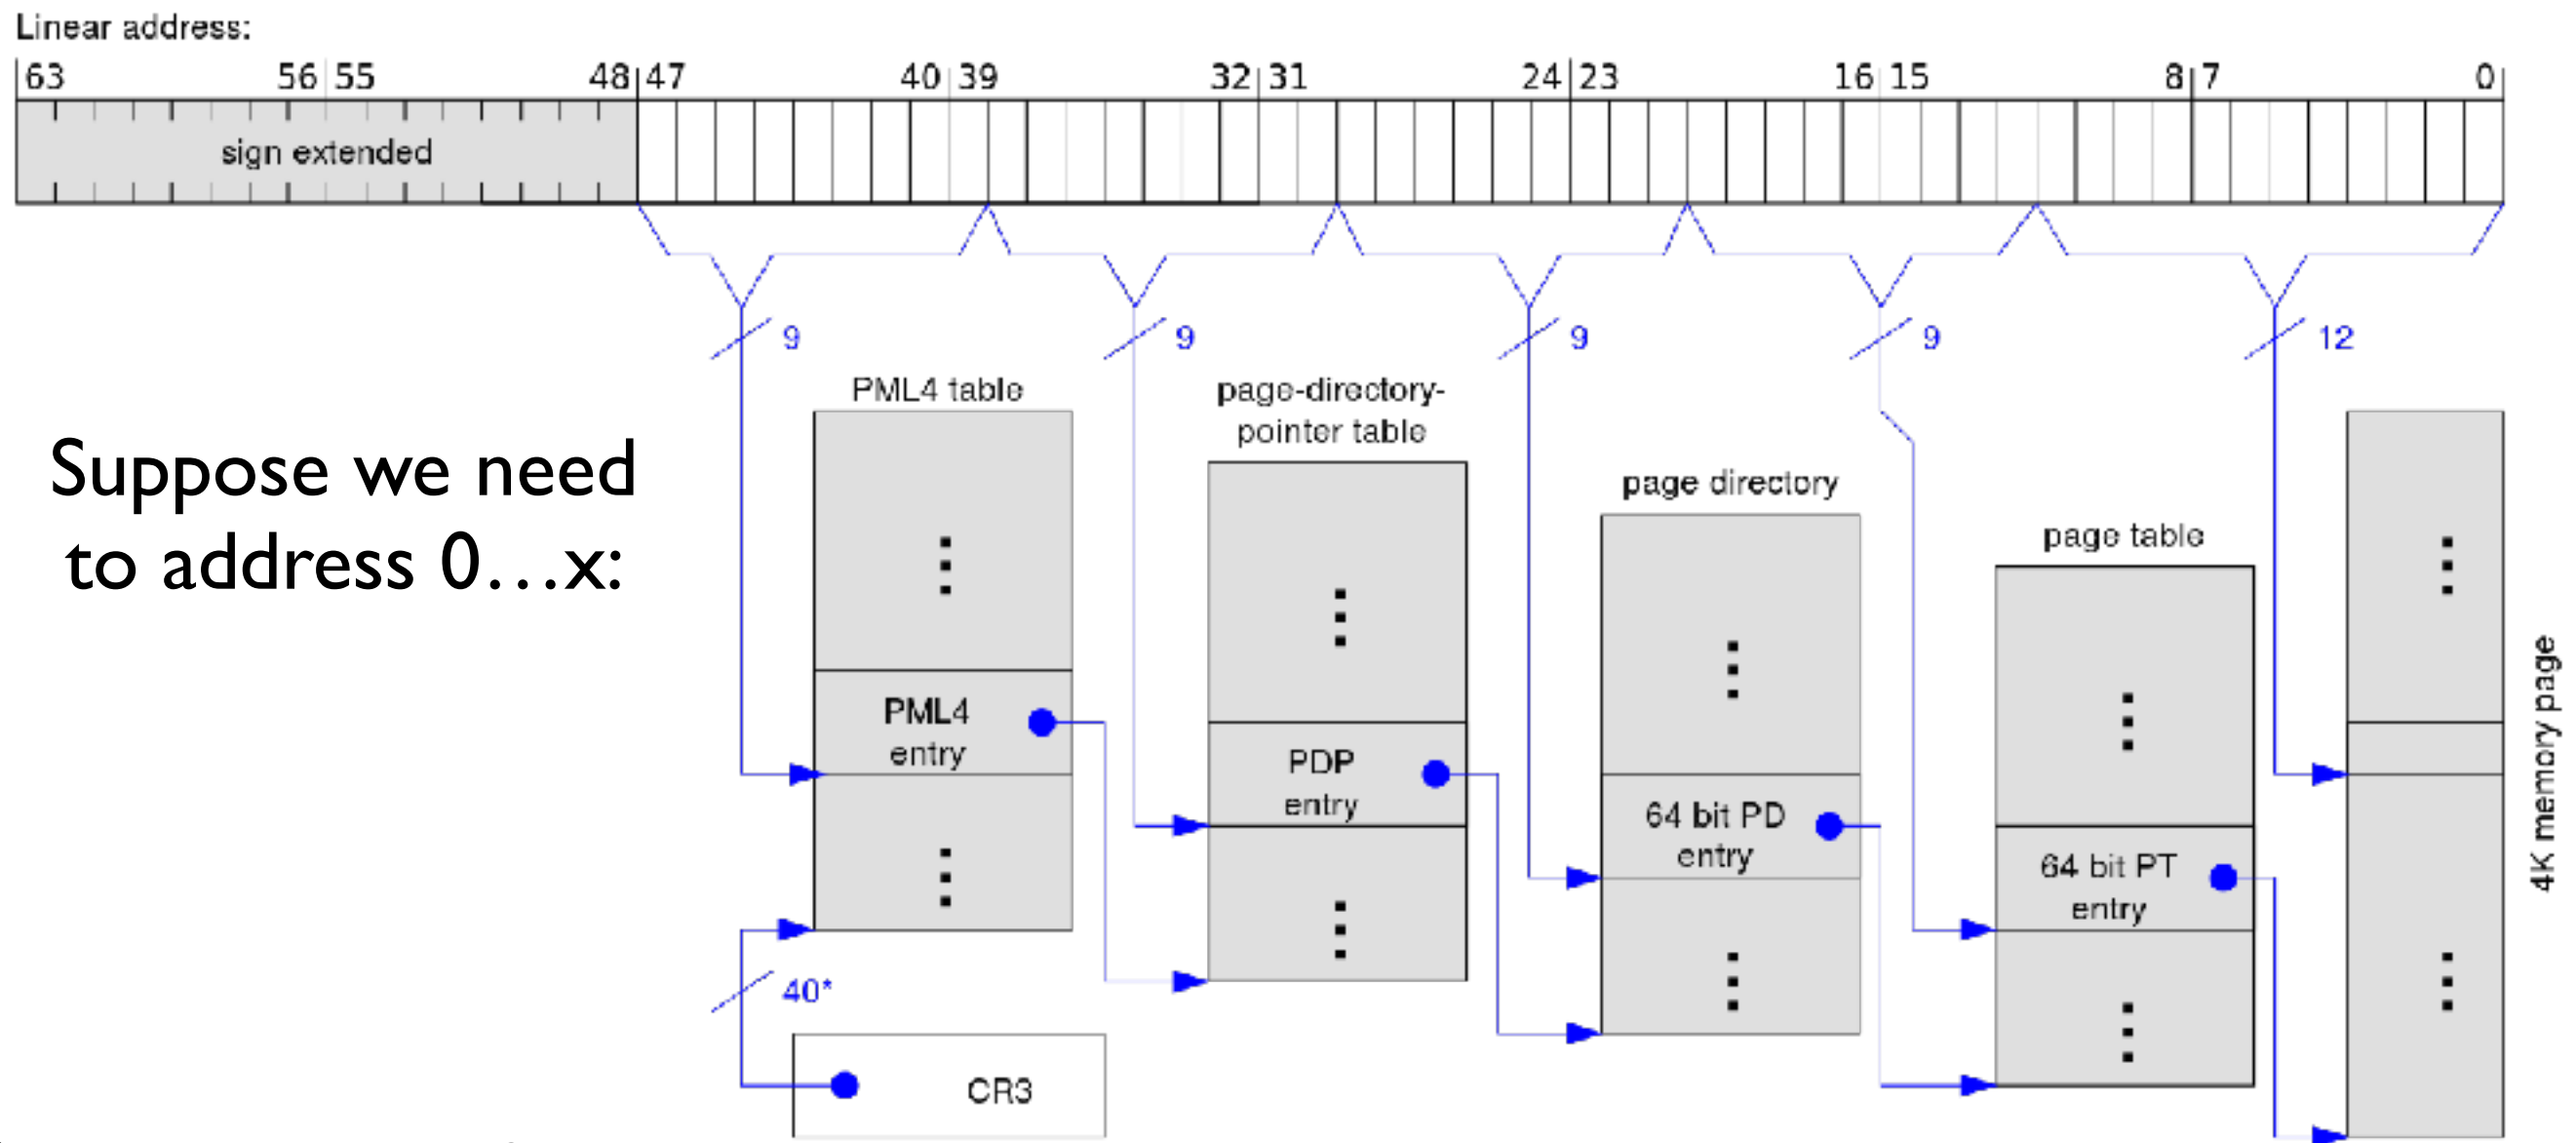
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

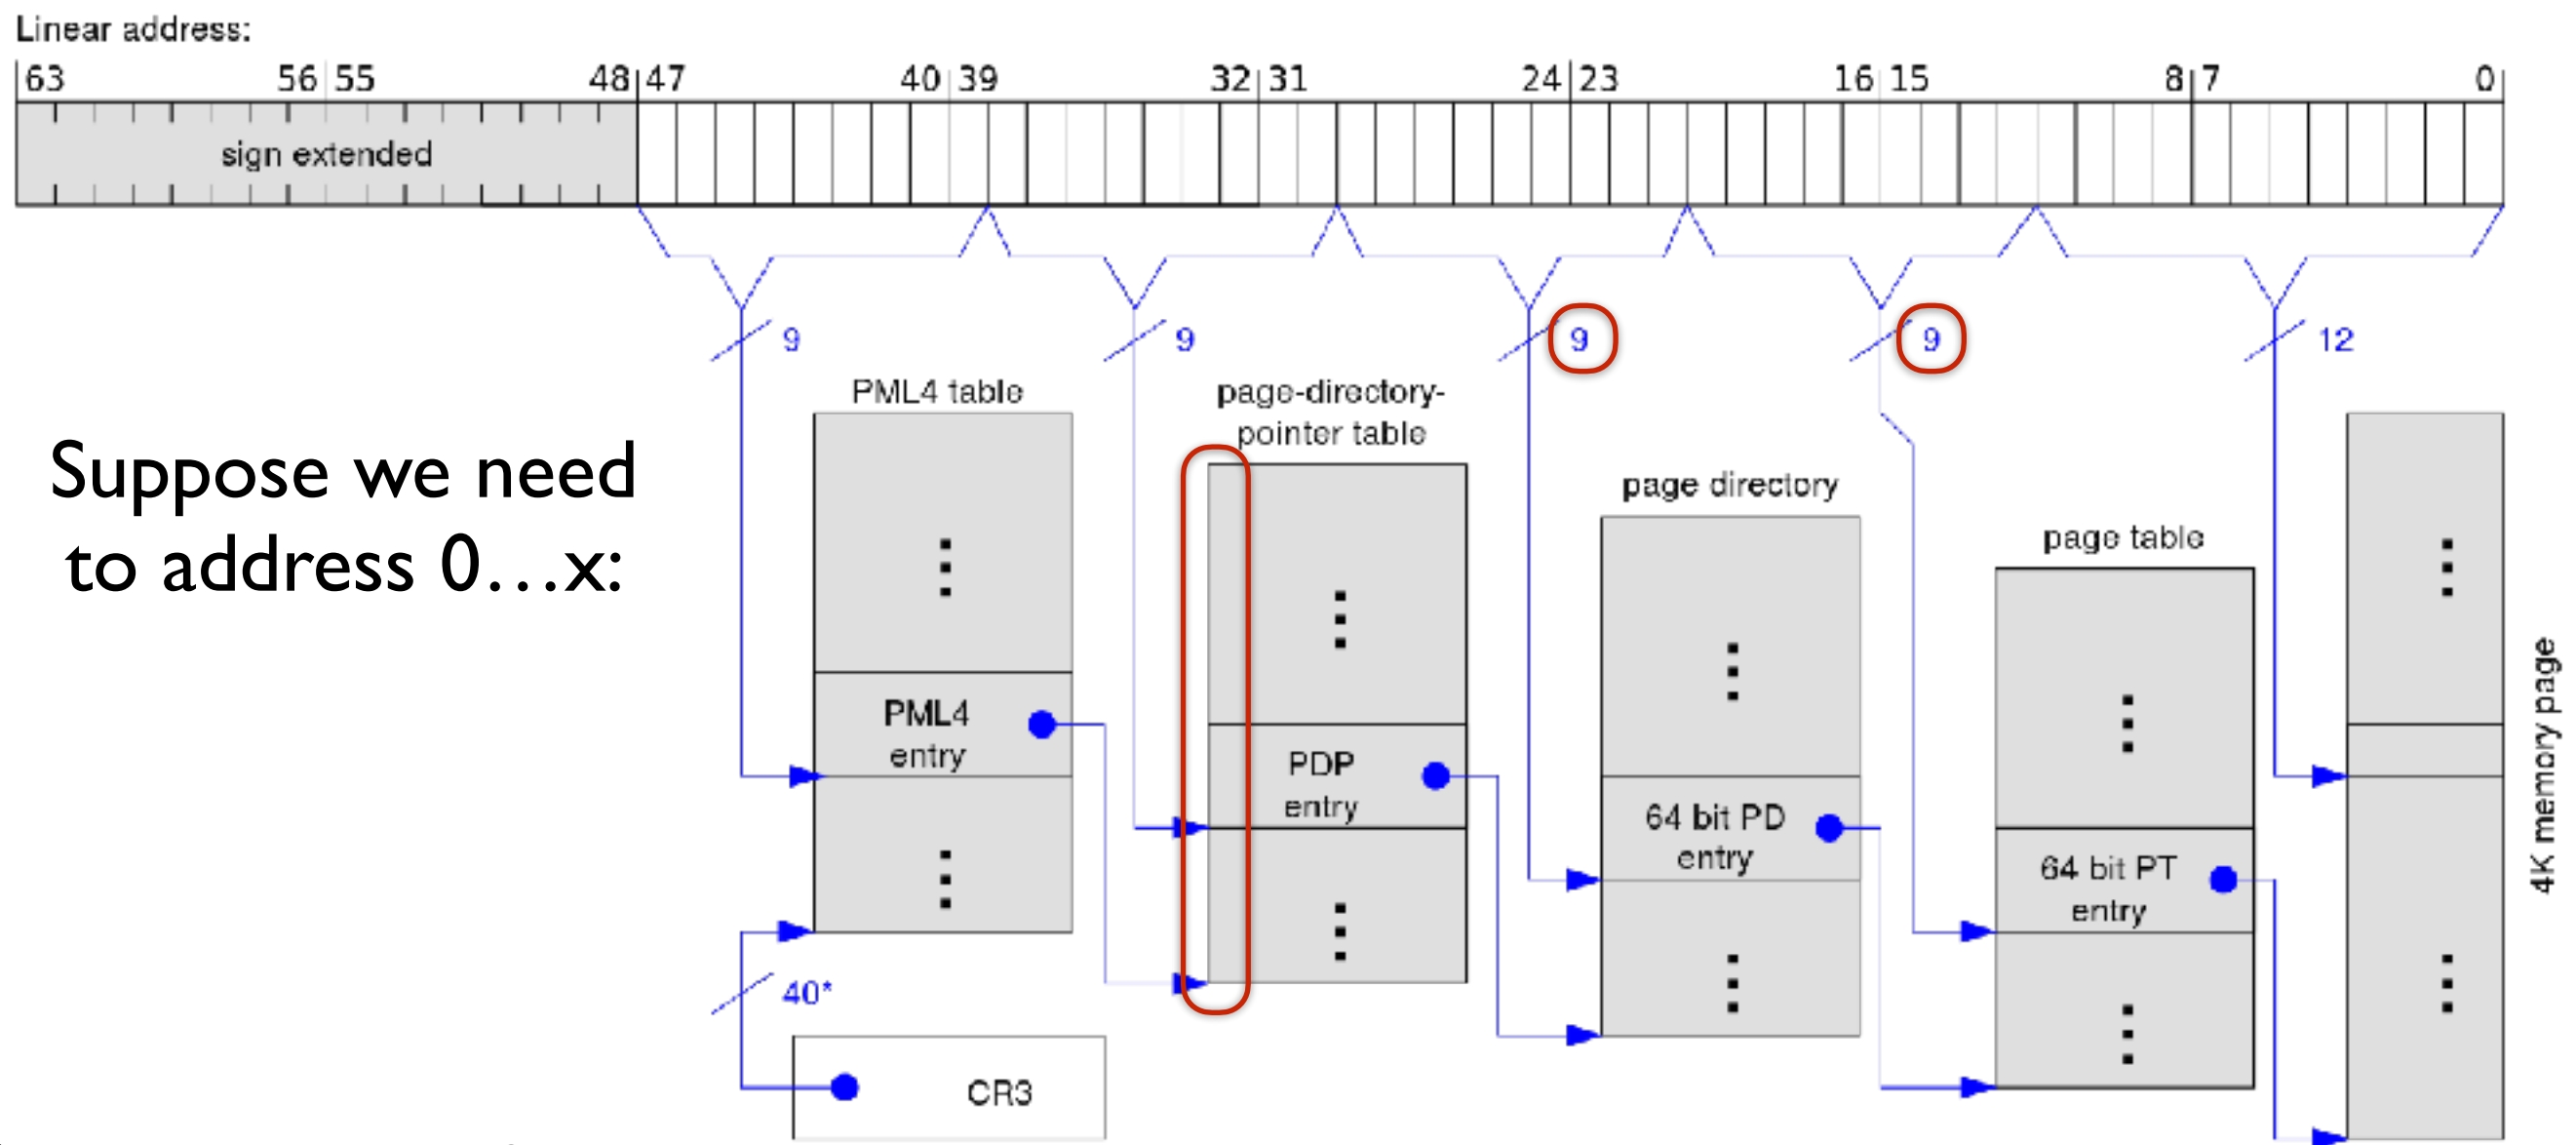
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

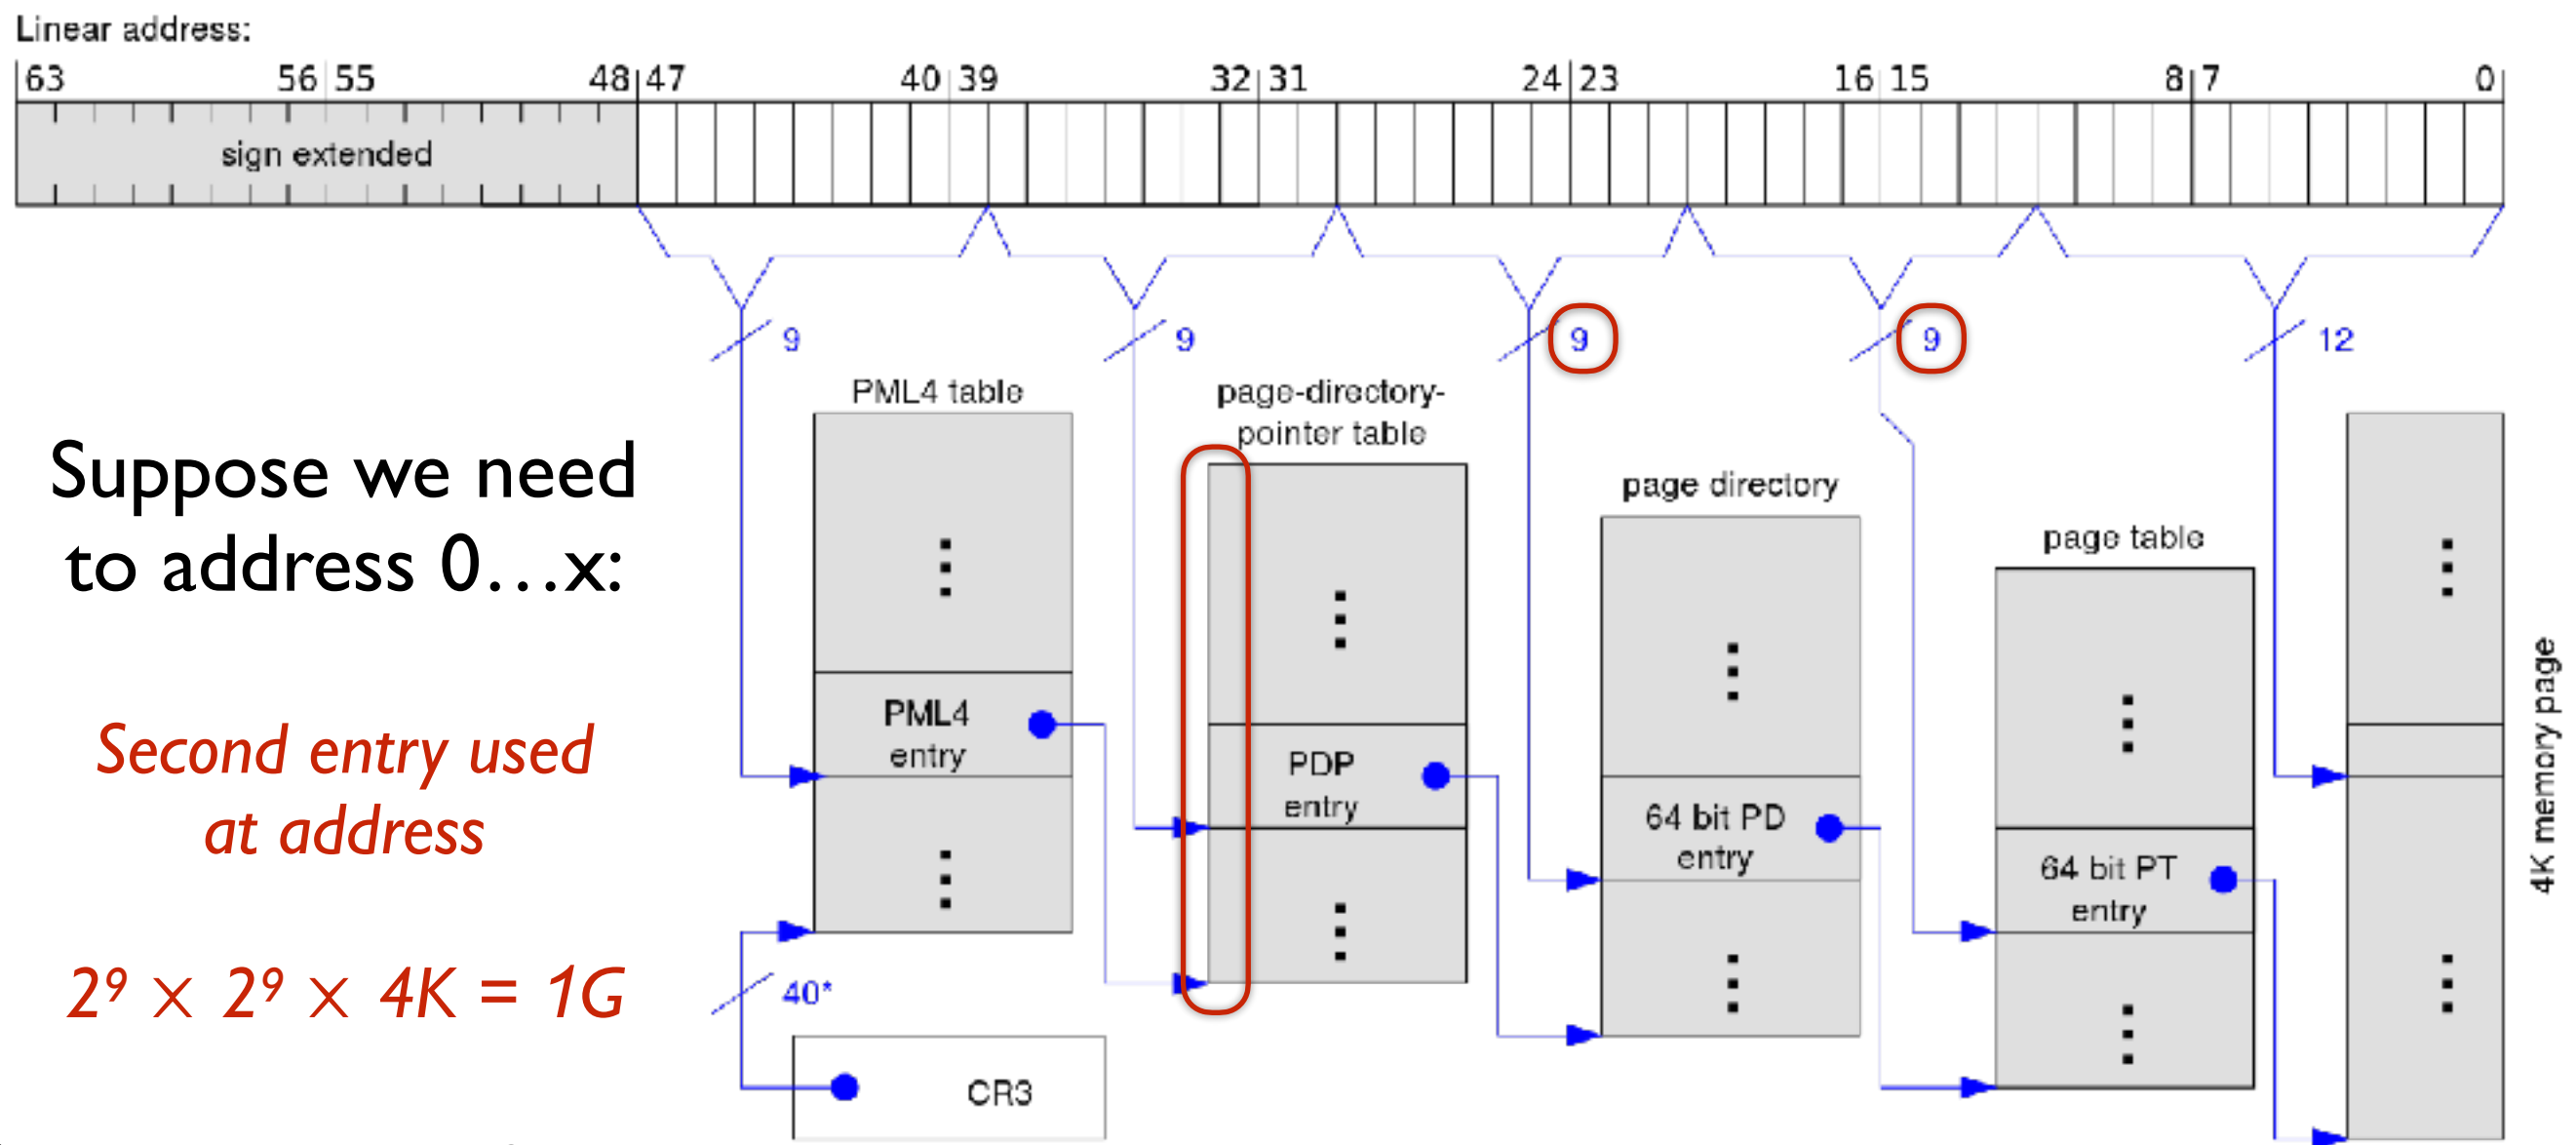
## Intel x86-64



\* source: Wikipedia user RokerHRO

# Multilevel Paging

## Intel x86-64



\* source: Wikipedia user RokerHRO

# Paging



# Paging

- Each process has its own page-table structure

# Paging

- Each process has its own page-table structure
- It's a *tree*, not a chain

# Paging

- Each process has its own page-table structure
- It's a *tree*, not a chain
  - May have looked like a chain in the previous slides...

# Paging

- Each process has its own page-table structure
- It's a *tree*, not a chain
  - May have looked like a chain in the previous slides...
  - But it's not! As your memory footprint increases, the more the structure looks like a tree

# Paging

- Each process has its own page-table structure
- It's a *tree*, not a chain
  - May have looked like a chain in the previous slides...
  - But it's not! As your memory footprint increases, the more the structure looks like a tree
- On context-switch, change it completely

# Paging

- Each process has its own page-table structure
- It's a *tree*, not a chain
  - May have looked like a chain in the previous slides...
  - But it's not! As your memory footprint increases, the more the structure looks like a tree
- On context-switch, change it completely
  - On x86-64, that means reassigning the CR3 register

# Paging

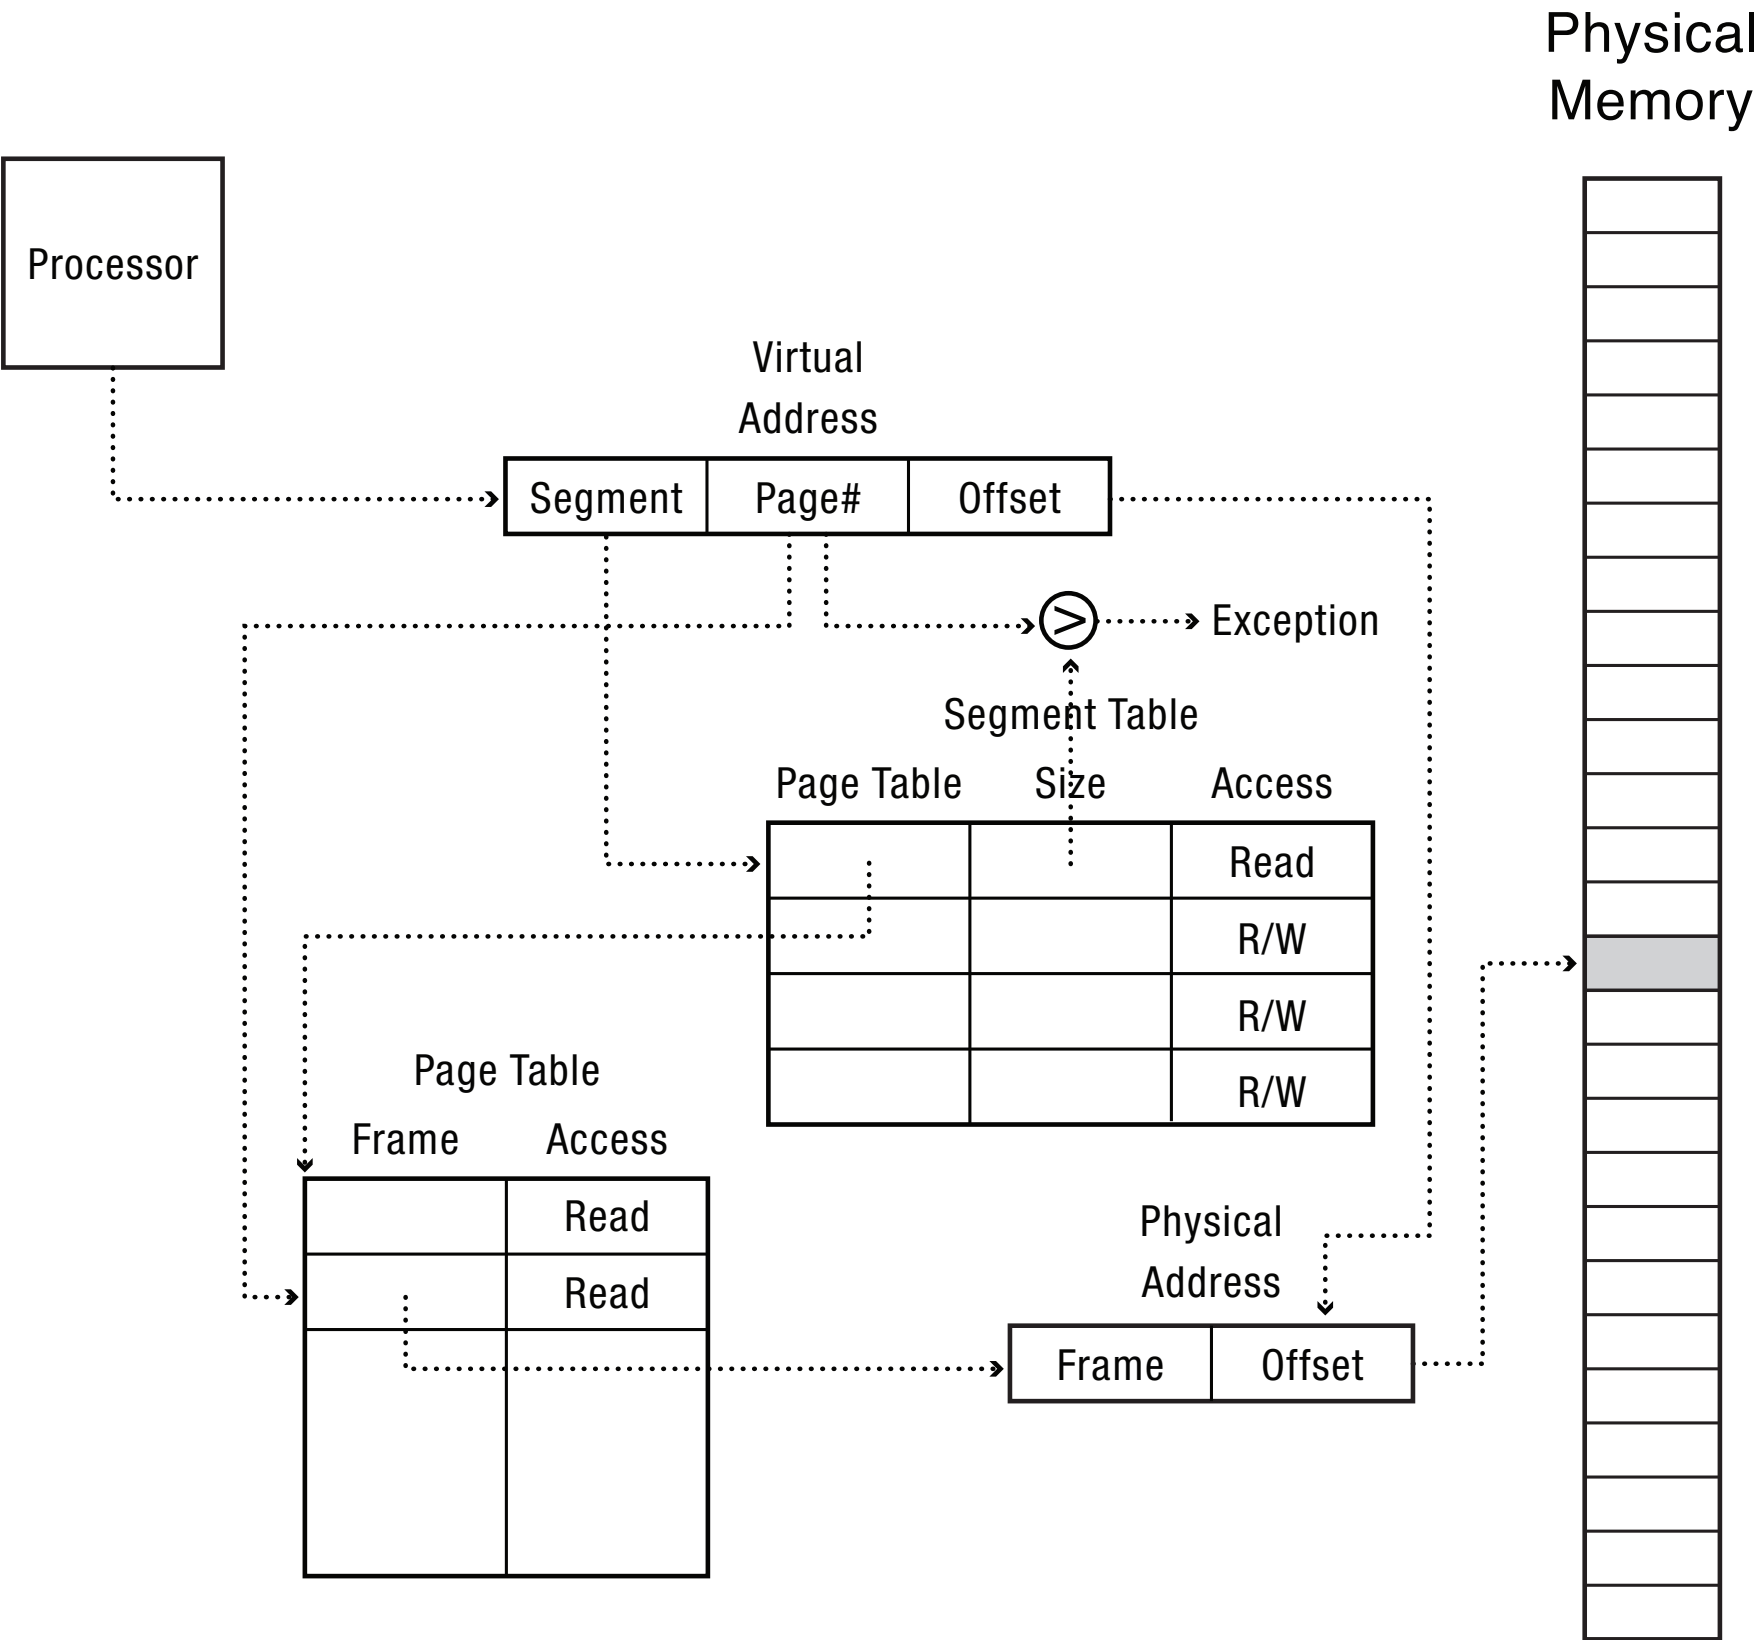
- Each process has its own page-table structure
- It's a *tree*, not a chain
  - May have looked like a chain in the previous slides...
  - But it's not! As your memory footprint increases, the more the structure looks like a tree
- On context-switch, change it completely
  - On x86-64, that means reassigning the CR3 register
  - One-shot change

# Alternative

(less important)



# Paged Segmentation



# Fim

\*all images belong to their copyright owners, including  
the Computer History Museum, Wikipedia/Youtube users, or Recursive Books